

PPE Detection Using Machine Learning For Enhanced Safety in Construction

Darin Verduzco, Fatimat Atanda, Victor Hugo Germano

Group 4

Shiley-Marcos School of Engineering, University of San Diego

AAI-501: Introduction to Artificial Intelligence

Andrew Van Benschoten, PhD

December 9, 2024

Abstract

Construction work professionals are subject to multiple risks in the work environment that are mitigated using protection, and relying on safety procedures to ensure usage. Using computer vision technologies we created a model capable of identifying people wearing Personal Protective Equipment (PPE). The model leverages a computer vision model, YOLOv8 (You Only Look Once), convolutional neural network and a dataset of images to train a model using two different approaches; custom manual-labelled approach and hybrid approach (combining manually annotated and auto-annotated images).

Introduction

Construction practitioners are at a disproportionately higher risk of fatal and nonfatal injuries compared to practitioners from other industries. The absence of and inappropriate use of personal protective equipment (PPE), hereinafter referred to as PPE non-compliance, are major causes of fatal and nonfatal injuries at construction workplaces (Jalil Al-Bayati et al., 2023). One strategy to mitigate these incidents is implementing a strict policy of wearing personal protective equipment (PPE) when working in hazardous environments. This equipment can include protective glasses, high-visibility vests, gloves, and helmets to ensure the workers remain safe and productive in the long run.

However, manually monitoring PPE compliance is challenging as it involves the daily supervision of an extensive workforce. Traditional safety inspections rely heavily on human intervention, which is prone to error and delays (Kursunoglu et al., 2022).

An alternative is using the latest computer vision (CV)-based solutions and frameworks to determine whether workers wear PPE according to safety protocols. By integrating computer vision PPE detection technology, industries can automate and enhance the compliance monitoring process. (Vukicevic et al., 2024).

Objectives

- Develop a YOLOv8-based object detection model for PPE compliance monitoring.
- Combine team manually annotated datasets with labelled Kaggle datasets for model training.
- Combine manually annotated and auto-annotated datasets for model training (hybrid approach).
- Evaluate the system's performance on a diverse dataset to ensure reliability.

Methodology

1. Data Collection and Preparation

Manually-annotated datasets included images created and labeled by the team to ensure high-quality bounding boxes and class annotations for PPE object detection. This dataset was used for the custom model training approach. It serves as a gold standard for initial training, providing a reliable foundation for the model to learn essential detection patterns in the combined dataset used for hybrid training. This ensures accuracy, eliminates labeling noise, and provides clear, consistent annotations for critical PPE categories.

The dataset for the custom model includes 2,801 images and labels from a Kaggle “Construction Site Safety Image Dataset Roboflow” dataset and 58 images extracted from 3 personal videos filmed by Darin Verduzco which were manually annotated using a free open

source labeling software, Computer Vision Annotation Tool (CVAT), by first defining the 10 class labels in the software, then drawing bounding boxes around each object in the images. There was no confirmation found that all annotations in the Kaggle dataset were created manually, such as in the images from our personal video source, so “initial annotation count” is the terminology used for annotations in Table 1.

Auto-annotated datasets used pre-trained YOLOv8 model to automatically label datasets. Annotations include bounding boxes and class labels for PPE items, with a confidence threshold of 0.5 applied to filter predictions. This approach allows for rapid scale up of dataset size, reducing the time and cost of manual labeling. It introduces variability in data to improve the model’s generalization capability.

Combined Dataset: The manually annotated and auto-annotated datasets were merged to form the final training dataset for the hybrid training approach. This datasets includes 63 images sourced from google in addition to custom datasets. Of the total datasets, 178 images were manually annotated and the rest of the images (2,746) were annotated using a pre-trained model for efficiency and scalability.

Dataset Labels/Classes: The 10 class labels from this dataset include the following: “Hardhat”, “Mask”, “NO-Hardhat”, “NO-Mask”, “NO-Safety Vest”, “Person”, “Safety Cone”, “Safety Vest”, “Machinery”, and “Vehicle”. The inclusion of “NO-” before the objects was desirable because it will be useful for providing counts when an object is not detected on a person. However, the focus is on 7 main objects identification; “Hardhat”, “NO-Hardhat”, “Mask”, “NO-Mask”, “Safety Vest”, “NO-Safety Vest”, “Person,”.

Data Combination: The custom datasets were randomly split into training (70%), validation (20%), and test (10%) subsets. For the hybrid datasets, the manually annotated datasets (178)

were split into training (80%), validation (20%) subsets while the final combined datasets were split into training (70%), validation (20%), and test (10%).

Dataset Properties: The personal video source and Kaggle dataset images contained image resolutions of 640 x 640 pixels. Image sizes of the google search images varied, but a conversion to the size of 640 x 640 will be done automatically by default when training an image dataset on a YOLO model. 27 filenames from the Kaggle dataset larger than 100 characters were also truncated manually to eliminate errors.

Dataset Description	Trained from Custom Dataset		Hybrid Model	
	Image Count	Initial Annotation Count	Image Count	Initial Annotation Count
Kaggle dataset: "Construction Site Safety Image Dataset Roboflow"	2801	2801	2801	120
Personal video source - Darin Verduzco	58	58	58	58
Google search images	0	0	63	0
Total	2859	2859	2922	178

Table 1: The summary of counts for the dataset of each model

2. Model Training

Custom Training Parameters

Model training was conducted on all annotated datasets (kaggle and custom) with the following parameters and lasted approximately 8.78 hours: epochs: 30, seed: 88 (for reproducibility), nms: True (Non-Maximum Suppression for predictions), optimizer : AdamW (via auto), and batch: 0.7 (determines batch size based on 70% memory usage). Prediction parameters used included iou: 0.5 (Intersection over Union threshold, for reducing duplicates) and conf: 0.4 (confidence threshold to remove low confidence predictions).

Hybrid Training Parameters

Initial training was conducted with manually annotated datasets first. The pretrained model was then used to auto-annotate the rest of the datasets by generating bounding boxes and class labels, as described in Table 2, leveraging prior knowledge from related datasets to bootstrap the dataset. Reducing the dependence on manual annotations saves time and resources while maintaining reasonably good annotation quality. The final training was then conducted on combined datasets (manual and auto-annotated).

Parameter	Initial Training	Final Training
Dataset	Manually Annotated Dataset	Combined Dataset
Epochs	50	30
Image Size	640	640
Batch Size	16	20
Learning Rate(lr_0)	1e-4	1e-3
Optimizer	AdamW(via “auto”)	AdamW(via “auto”)
Mixed Precision(amp)	Not Used	Enabled (for faster training)
Patience (early stopping)	10	10

Table 2: Hybrid Model Training Parameters

Results

Evaluation Metrics

- **Precision (P):** Measures how many of the model’s predictions are correct.
- **Recall (R):** Measures how many true instances the model captures.
- **Mean Average Precision (mAP50):** Evaluates detection accuracy at an IoU threshold of 50%.
- **Mean Average Precision (mAP50-95):** Evaluates detection accuracy across multiple IoU thresholds.

The tables 3 and 4 show the class-specific and overall model performance on custom dataset training and hybrid dataset training.

Class	Precision	Recall	mAP50	mAP50-95
Mask	0.933	0.911	0.94	0.743
Safety Vest	0.918	0.759	0.834	0.601
Hardhat	0.918	0.724	0.807	0.529
Person	0.923	0.812	0.892	0.681
NO-Mask	0.875	0.657	0.722	0.399
NO-Safety Vest	0.908	0.762	0.834	0.588
NO-Hardhat	0.93	0.702	0.806	0.529
Overall (10 classes)	0.904	0.752	0.824	0.581

Table 3: Custom Datasets Results

Class	Precision	Recall	mAP50	mAP50-95
Mask	0.829	0.946	0.929	0.782
Safety Vest	0.785	0.842	0.89	0.709
Hardhat	0.863	0.956	0.972	0.83
Person	0.885	0.92	0.956	0.825
NO-Mask	0.772	0.83	0.841	0.56
NO-Safety Vest	0.722	0.73	0.808	0.62
NO-Hardhat	0.833	0.875	0.915	0.737
Overall (10 classes)	0.805	0.868	0.89	0.715

Table 4: Hybrid Datasets Results

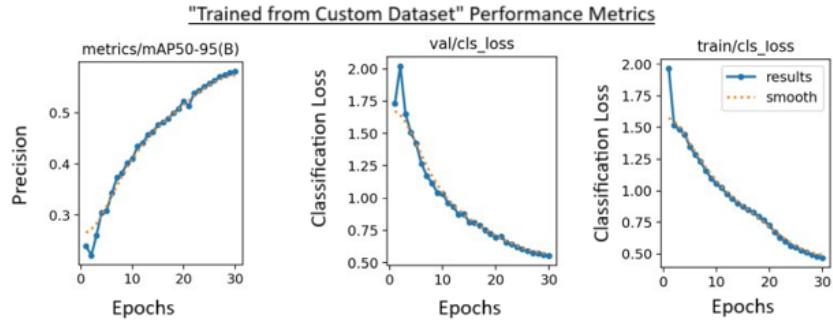


Figure 1: Custom Dataset Performance Metrics

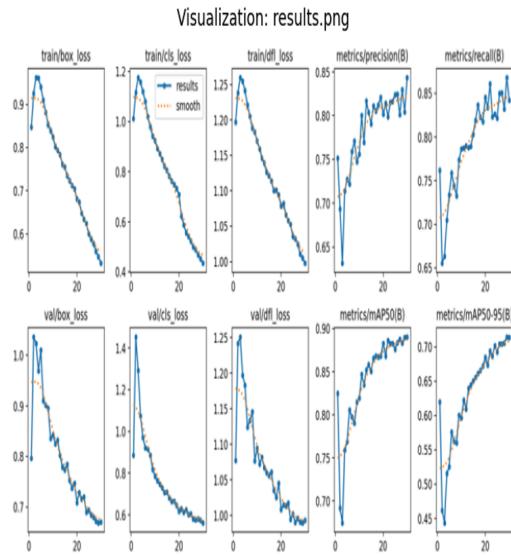


Figure 2: Hybrid Training Performance Metric

Analysis

The model demonstrated consistent improvement across epochs in both training approaches, with an increasing mAP50 score and steadily decreasing losses (box loss, class loss, and DFL loss), indicating good convergence. High precision and recall values highlight the model's reliability, while validation images confirm its ability to detect multiple objects in complex construction scenes. Additionally, the model showed strong generalization to unseen images, suggesting its potential for real-world construction site safety applications. However, the significant gap

between mAP50 and mAP50-95 scores indicates a decline in performance at higher IoU thresholds. While the model performed well across most classes, it struggled slightly with "NO-Mask" detection presented in Table 4, revealing room for improvement. Some misclassifications and missed detections were observed, which could be mitigated by addressing class imbalances, increasing the dataset size for underperforming classes, and employing techniques such as transfer learning or ensemble methods.

Real Life Example



Figure 3 - FUTEK HD Modern Basement Construction Technology

Testing the model on a real construction site's video stream provides insight into its readiness for production. Figure 3 highlights key challenges such as image quality, camera distance, non-workers on-site, and non-fixed cameras. Addressing these issues involves establishing standard operations for video capture and image processing, including fixed camera positions and geofencing to exclude non-site entities (e.g., spectators, vehicles, animals) and reduce false positives.

The next step is enhancing the model's categorization capabilities through a combination of refinement techniques, performance optimization, and bias mitigation. The first step is model and dataset refinement: fine tuning the hyperparameters for training, exploring different architectures and amplifying the use of augmentation techniques. Balancing the class distribution and expanding the dataset, reducing the risk of model bias. Techniques like Layer-wise Relevance Propagation (LRP) can be used to minimize the influence of background bias on classifiers.

Performance Optimization could play a significant role. Applying techniques like filter pruning or low-rank approximation to reduce model size without significantly impacting performance and investigating bias in this dataset can also prove to be a good decision: Using methods like Hybrid Sample Synthesis to identify likely bias-conflicting samples and creating new training data to reduce the impact of bias-aligned samples.

Conclusion

The development of a PPE detection model using the YOLOv8 framework demonstrated high precision and recall in detecting critical PPE categories, showcasing its potential for improving workplace safety. The hybrid training approach effectively combined manually and auto-annotated datasets, achieving robust detection performance.

Both models performed to a satisfactory degree, with the hybrid approach having a much larger training time performance. As the datasets increase in the number of images not annotated, hybrid would be the preferred option, understanding the possibility of bias and unbalanced classes represented in the dataset.

Key recommendations for improvement include expanding the dataset to enhance generalizability, adopting active learning for continuous refinement, prioritizing real-time deployment with alert systems, and implementing robust performance monitoring during deployment. These enhancements aim to ensure sustained accuracy, scalability, and adaptability in diverse environments.

The project underscores the transformative potential of automated PPE detection systems in mitigating workplace risks, supporting safety compliance, and promoting safer work environments. It highlights the intersection of technology and occupational safety, advancing automation in compliance management while safeguarding workers.

References

- Computer Vision Annotation Tool. (2024). Accessed November 25, 2024, from <https://www.cvat.ai/>
- Hady, M. F. A., & Schwenker, F. (2013). Semi-supervised learning. In M. Bianchini, M. Maggini, & L. Jain (Eds.), *Handbook on neural information processing* (Vol. 49, pp. 215\u2013244). Springer. https://doi.org/10.1007/978-3-642-36657-4_7
- Jalil Al-Bayati, A., Rener, A. T., Listello, M. P., & Mohamed, M. (2023). PPE non-compliance among construction workers: An assessment of contributing factors utilizing fuzzy theory. In *Journal of Safety Research* (85), 242–253. <https://doi.org/10.1016/j.jsr.2023.02.008>
- Sanyal, S. (2022). Construction Site Safety Image Dataset Roboflow, Version 1. Retrieved November 2, 2024 from <https://www.kaggle.com/datasets/snehilsanyal/construction-site-safety-image-dataset-roboflow/data>
- OpenAI. (2024). *ChatGPT* (v2.0) [Large Language Model]. <https://chat.openai.com/chat>
- Ultralytics. (2024). *YOLO performance metrics: A guide*. Retrieved December 1, 2024, from <https://docs.ultralytics.com/guides/yolo-performance-metrics/>
- Verduzco, D. (2024) Construction PPE Personal Videos, OneDrive. Retrieved November 25, 2024 from https://uofsandiego-my.sharepoint.com/:f/g/personal/dverduzco_sandiego_edu/EpWCwg-E6jZGltY-j_EcIUkBIgGk1fMxfnBQKqBTSfaVTg?e=42Zkn7
- *Workers at construction sites wearing PPE*. (n.d.). Retrieved November 25, 2024, from Google Images: <https://www.google.com>

Train from Custom Dataset

December 8, 2024

1 Construction PPE Object Detection:

1.0.1 Train From Custom Dataset

AAI-501 Group 4: Fatimat Atanda, Victor Hugo Germano, Darin Verduzco GitHub:
<https://github.com/victorhg/aai-501-final-project>

2 Dataset setup:

download custom “datasets” folder into “train_from_custom” folder (will be ignored to git).
Dataset Summary: construction_ppe: 2801 images custom dataset: 58 images totals: (2859 raw_images and 2859 raw_labels) (note: file names larger than 99 characters may not be read, update manually)

2.0.1 Import Modules and Libraries

```
[5]: import os
import sys
import cv2 # OpenCV for images
import time
import glob
import yaml
import torch
import shutil
import random
import warnings
import numpy as np
import pandas as pd
from PIL import Image
from pathlib import Path
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import albumentations as A # For image augmentation
from ultralytics import YOLO # For object detection
from albumentations.pytorch import ToTensorV2 # For image formating
from torchvision.transforms import ColorJitter
```

```
[6]: warnings.filterwarnings("ignore") # Suppress all warnings
# Import msaaai from parent folder
module_path = os.path.abspath(os.path.join('..'))
if module_path not in sys.path:
    sys.path.append(module_path+"/")
import msaaai
```

2.0.2 Set path directories

```
[5]: dataset_base_dir = "datasets"
image_base_dir = dataset_base_dir + "/images"
labels_base_dir = dataset_base_dir + "/labels"

# Define directories to reset
image_dirs = {
    "train": f"{image_base_dir}/train",
    "valid": f"{image_base_dir}/valid",
    "test": f"{image_base_dir}/test",
}
label_dirs = {
    "train": f"{labels_base_dir}/train",
    "valid": f"{labels_base_dir}/valid",
    "test": f"{labels_base_dir}/test",
}
```

2.0.3 Custom_data.yaml file creation in dataset dir

```
[31]: ppe_classes = [
    'Hardhat',
    'Mask',
    'NO-Hardhat',
    'NO-Mask',
    'NO-Safety Vest',
    'Person',
    'Safety Cone',
    'Safety Vest',
    'Machinery',
    'Vehicle'
]
number_classes = len(ppe_classes)
output_dir = 'datasets'

dict_file = {
    'train': 'images/train',
    'val': 'images/valid',
    'test': 'images/test',
    'nc': number_classes,
```

```

    'names': ppe_classes
}

with open(os.path.join(dataset_base_dir, 'custom_data.yaml'), 'w+') as file:
    yaml.dump(dict_file, file)
# Full yaml path to prevent path error
yaml_path = str(Path('datasets/custom_data.yaml').resolve())

```

2.0.4 Clear train/test/valid image folders from “datasets” folder

```
[23]: directories_to_reset = list(image_dirs.values()) + list(label_dirs.values())

# Reset directories
msaai.reset_directories(directories_to_reset)

print("Directories reset and ready for use.")
```

Deleted existing directory: datasets/images/train
 Recreated directory: datasets/images/train
 Deleted existing directory: datasets/images/valid
 Recreated directory: datasets/images/valid
 Deleted existing directory: datasets/images/test
 Recreated directory: datasets/images/test
 Deleted existing directory: datasets/labels/train
 Recreated directory: datasets/labels/train
 Deleted existing directory: datasets/labels/valid
 Recreated directory: datasets/labels/valid
 Deleted existing directory: datasets/labels/test
 Recreated directory: datasets/labels/test
 Directories reset and ready for use.

2.0.5 Split into train/valid/test folders from “raw_image” and “raw_label” folders

```
[24]: source_image_dir = dataset_base_dir + "/raw_images"
source_label_dir = dataset_base_dir + "/raw_labels"

# Combine and assign to DataFrame
img_split_df = msaai.img_train_test_split(source_image_dir)

# Copy files from source to folders
msaai.arrange_image_and_label_files(
    img_split_df,
    source_image_dir,
    source_label_dir,
    image_dirs,
    label_dirs
)
```

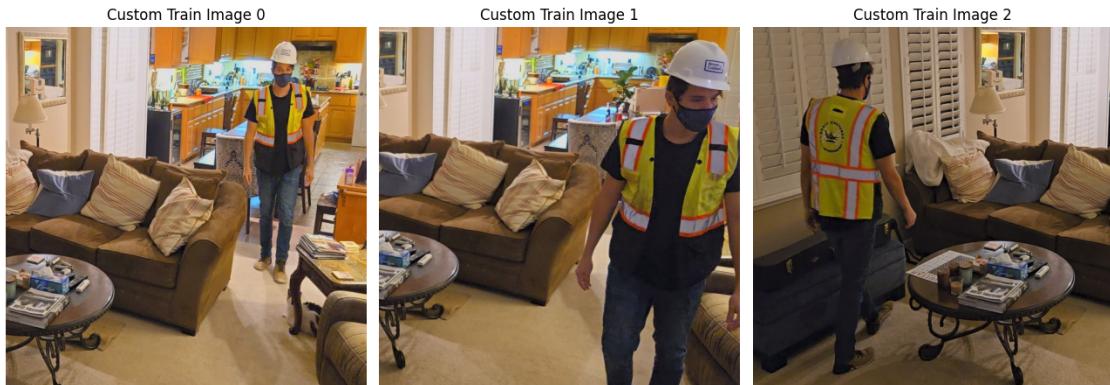
2.0.6 Show custom image dataset

```
[22]: # Get the paths to the first three custom images
image_paths = glob.glob(f"[image_dirs['train']]*/scene*.jpg")[:3]

# Check if there are enough images
if len(image_paths) < 3:
    print("Not enough images found!")
else:
    # Create a subplot with 1 row and 3 columns
    fig, axes = plt.subplots(1, 3, figsize=(13, 7))  # Width, height

    for i, ax in enumerate(axes):
        # Load and display each image
        image = mpimg.imread(image_paths[i])
        ax.imshow(image)
        ax.axis('off')  # Hide axes
        ax.set_title(f"Custom Train Image {i}")

    plt.tight_layout()
    plt.show()
```



2.0.7 Show construction dataset

```
[23]: # Get the paths to the first three construction images
image_paths = glob.glob(f"[image_dirs['train']]*/construction*.jpg")[:3]

# Check if there are enough images
if len(image_paths) < 3:
    print("Not enough images found!")
else:
    # Create a subplot with 1 row and 3 columns
```

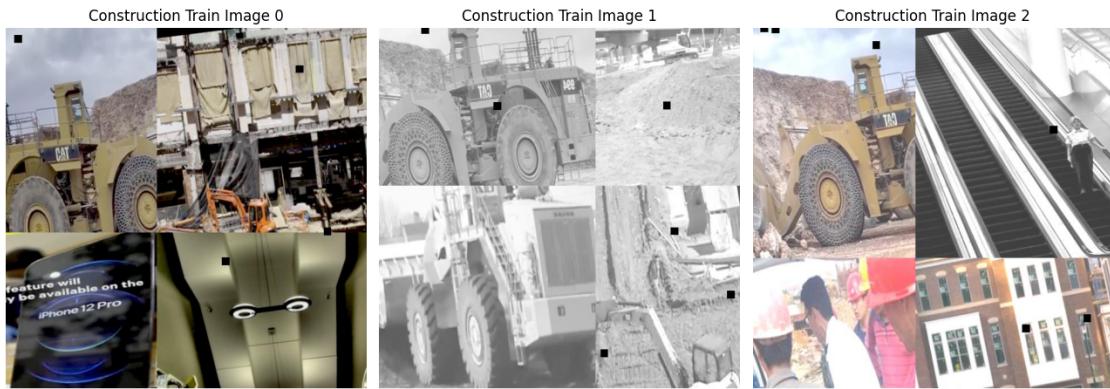
```

fig, axes = plt.subplots(1, 3, figsize=(13, 7)) # Adjust the figsize as needed

for i, ax in enumerate(axes):
    # Load and display each image
    image = mpimg.imread(image_paths[i])
    ax.imshow(image)
    ax.axis('off') # Hide axes
    ax.set_title(f"Construction Train Image {i}")

plt.tight_layout()
plt.show()

```



2.1 Train YOLOv8 from construction and custom dataset

```

[32]: proj_name = input('Please enter run folder description: ')
# Count in minutes
start = time.time()
for i in range(10000000):
    pass

# Train
model = YOLO('..../models/yolov8m.pt')
model.train(
    data = yaml_path,
    epochs = 30, # Dataset iterations
    patience = 10, # Epochs to wait to discontinue after no val metric improvement
    seed = 88, # Repeatability
    plots = True, # produce plots
    nms = True, # Non-Maximum Suppression (NMS) improves post-processing detection
)

```

```

batch = 0.7, # Determine batch size using defined fraction of memory usage
name = proj_name # input for run description
)

# Time end
end = time.time()
elapsed_seconds = end - start
elapsed_minutes = elapsed_seconds / 60
print(f"Total model train time: {elapsed_minutes:.2f} minutes")

```

Please enter run folder description: new m lower batch

Ultralytics 8.3.32 Python-3.12.4 torch-2.5.1+cpu CPU (Intel Core(TM) i9-14900KF)

engine\trainer: task=detect, mode=train,
model=../../models/yolov8m.pt, data=C:\Users\DAReN\Documents\Python-JL\Github\aaI-501-final
project\notebooks\train_from_custom\datasets\custom_data.yaml, epochs=30,
time=None, patience=10, batch=0.7, imgsz=640, save=True, save_period=-1,
cache=False, device=None, workers=8, project=None, name=new m lower batch,
exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=88,
deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10,
resume=False, amp=True, fraction=1.0, profile=False, freeze=None,
multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True,
split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300,
half=False, dnn=False, plots=True, source=None, vid_stride=1,
stream_buffer=False, visualize=False, augment=False, agnostic_nms=False,
classes=None, retina_masks=False, embed=None, show=False, save_frames=False,
save_txt=False, save_conf=False, save_crop=False, show_labels=True,
show_conf=True, show_boxes=True, line_width=None, format=torchscript,
keras=False, optimize=False, int8=False, dynamic=False, simplify=True,
opset=None, workspace=4, nms=True, lr0=0.01, lrf=0.01, momentum=0.937,
weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1,
box=7.5, cls=0.5, df1=1.5, pose=12.0, kobj=1.0, label_smoothing=0.0, nbs=64,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5,
shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0,
mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment,
erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml,
save_dir=runs\detect\new m lower batch
Overriding model.yaml nc=80 with nc=10

	from	n	params	module
arguments				
0		-1 1	1392	ultralytics.nn.modules.conv.Conv
[3, 48, 3, 2]				
1		-1 1	41664	ultralytics.nn.modules.conv.Conv
[48, 96, 3, 2]				
2		-1 2	111360	ultralytics.nn.modules.block.C2f

```

[96, 96, 2, True]
 3           -1 1    166272 ultralytics.nn.modules.conv.Conv
[96, 192, 3, 2]
 4           -1 4    813312 ultralytics.nn.modules.block.C2f
[192, 192, 4, True]
 5           -1 1    664320 ultralytics.nn.modules.conv.Conv
[192, 384, 3, 2]
 6           -1 4   3248640 ultralytics.nn.modules.block.C2f
[384, 384, 4, True]
 7           -1 1   1991808 ultralytics.nn.modules.conv.Conv
[384, 576, 3, 2]
 8           -1 2   3985920 ultralytics.nn.modules.block.C2f
[576, 576, 2, True]
 9           -1 1   831168 ultralytics.nn.modules.block.SPPF
[576, 576, 5]
10          -1 1      0 torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']
11          [-1, 6] 1      0 ultralytics.nn.modules.conv.Concat
[1]
12          -1 2   1993728 ultralytics.nn.modules.block.C2f
[960, 384, 2]
13          -1 1      0 torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']
14          [-1, 4] 1      0 ultralytics.nn.modules.conv.Concat
[1]
15          -1 2   517632 ultralytics.nn.modules.block.C2f
[576, 192, 2]
16          -1 1   332160 ultralytics.nn.modules.conv.Conv
[192, 192, 3, 2]
17          [-1, 12] 1      0 ultralytics.nn.modules.conv.Concat
[1]
18          -1 2   1846272 ultralytics.nn.modules.block.C2f
[576, 384, 2]
19          -1 1   1327872 ultralytics.nn.modules.conv.Conv
[384, 384, 3, 2]
20          [-1, 9] 1      0 ultralytics.nn.modules.conv.Concat
[1]
21          -1 2   4207104 ultralytics.nn.modules.block.C2f
[960, 576, 2]
22          [15, 18, 21] 1   3781486 ultralytics.nn.modules.head.Detect
[10, [192, 384, 576]]
Model summary: 295 layers, 25,862,110 parameters, 25,862,094 gradients, 79.1
GFLOPs

```

Transferred 469/475 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs\detect\new m
lower batch', view at <http://localhost:6006/>
Freezing layer 'model.22.dfl.conv.weight'

```

train: Scanning C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\la
train: WARNING  C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\images\trai
n\004720.jpg.rf.afc486560a4004c7cf67910af31a29c.jpg: 1 duplicate labels removed
train: WARNING  C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\images\trai
n\construction-813-.jpg.rf.b085952261fd98f2e76b8065de149b5f.jpg: 1 duplicate
labels removed
train: New cache created: C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-
project\notebooks\train_from_custom\datasets\labels\train.cache
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01,
blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3,
method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0),
tile_grid_size=(8, 8))
AutoBatch: Computing optimal batch size for imgsz=640 at 70.0% CUDA
memory utilization.
AutoBatch: intended for CUDA devices, using default batch-size 16
train: Scanning C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\la
train: WARNING  C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\images\trai
n\004720.jpg.rf.afc486560a4004c7cf67910af31a29c.jpg: 1 duplicate labels removed
train: WARNING  C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\images\trai
n\construction-813-.jpg.rf.b085952261fd98f2e76b8065de149b5f.jpg: 1 duplicate
labels removed
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01,
blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3,
method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0),
tile_grid_size=(8, 8))

val: Scanning C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-project\notebooks\train_from_custom\datasets\labe
val: New cache created: C:\Users\DARiN\Documents\Python-
JL\Github\aaI-501-final-
project\notebooks\train_from_custom\datasets\labels\valid.cache

Plotting labels to runs\detect\new_m_lower_batch\labels.jpg...
optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and
'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum'
automatically...
optimizer: AdamW(lr=0.000714, momentum=0.9) with parameter groups

```

77 weight(decay=0.0), 84 weight(decay=0.0005), 83 bias(decay=0.0)
TensorBoard: model graph visualization added
 Image sizes 640 train, 640 val
 Using 0 dataloader workers
 Logging results to runs\detect\new m lower batch
 Starting training for 30 epochs...

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	1/30	0G	1.299	1.966	1.479	262	640:
		125/125 [16:54<00:00,					
			Class	Images	Instances	Box(P)	R
mAP50-95): 100%			18/18 [01:46				mAP50
				all	569	7670	0.448
0.24							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	2/30	0G	1.239	1.515	1.428	275	640:
		125/125 [16:04<00:00,					
			Class	Images	Instances	Box(P)	R
mAP50-95): 100%			18/18 [01:51				mAP50
				all	569	7670	0.436
0.221							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	3/30	0G	1.253	1.483	1.422	231	640:
		125/125 [16:20<00:00,					
			Class	Images	Instances	Box(P)	R
mAP50-95): 100%			18/18 [01:51				mAP50
				all	569	7670	0.487
0.26							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	4/30	0G	1.242	1.444	1.415	326	640:
		125/125 [15:45<00:00,					
			Class	Images	Instances	Box(P)	R
mAP50-95): 100%			18/18 [01:48				mAP50
				all	569	7670	0.537
0.304							

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
5/30 100%	0G 125/125 [15:11<00:00, mAP50-95): 100%	1.202 Class 18/18 [01:44 all 0.308	1.346 Images 18/18 [01:44 569	1.382 Instances 7670	331 Box(P R 0.501	640: mAP50 0.539

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
6/30 100%	0G 125/125 [15:18<00:00, mAP50-95): 100%	1.173 Class 18/18 [01:45 all 0.343	1.287 Images 18/18 [01:45 569	1.366 Instances 7670	275 Box(P R 0.522	640: mAP50 0.592

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
7/30 100%	0G 125/125 [15:12<00:00, mAP50-95): 100%	1.15 Class 18/18 [01:46 all 0.373	1.233 Images 18/18 [01:46 569	1.349 Instances 7670	255 Box(P R 0.566	640: mAP50 0.636

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
8/30 100%	0G 125/125 [15:16<00:00, mAP50-95): 100%	1.106 Class 18/18 [01:49 all 0.382	1.158 Images 18/18 [01:49 569	1.327 Instances 7670	325 Box(P R 0.571	640: mAP50 0.65

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	9/30	0G	1.089	1.097	1.304	244	640:
100%	125/125 [16:22<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:56	all	569	7670	0.782	0.594	0.668
			0.402				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
10/30	0G	1.067	1.056	1.283	424	640:	
100%	125/125 [17:21<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:50	all	569	7670	0.805	0.591	0.679
	0.411						

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
11/30	0G	1.037	1.023	1.267	274	640:	
100%	125/125 [16:16<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:49	all	569	7670	0.818	0.622	0.705
	0.434						

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
12/30	0G	1.022	0.9724	1.257	376	640:	
100%	125/125 [16:11<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:51	all	569	7670	0.828	0.63	0.712
	0.441						

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
13/30	0G	0.9995	0.9381	1.239	556	640:	
100%	125/125 [16:14<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:52	all	569	7670	0.828	0.63	0.712

	all	569	7670	0.827	0.64	0.723
0.457						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
14/30	0G	0.9771	0.899	1.225	276	640:
100%	125/125 [16:13<00:00,	Class	Images	Instances	Box(P	R
mAP50-95): 100%	18/18 [01:51					mAP50
	all	569	7670	0.828	0.652	0.73
0.462						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
15/30	0G	0.9546	0.8733	1.216	308	640:
100%	125/125 [16:33<00:00,	Class	Images	Instances	Box(P	R
mAP50-95): 100%	18/18 [01:54					mAP50
	all	569	7670	0.85	0.666	0.744
0.476						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
16/30	0G	0.9424	0.8487	1.203	308	640:
100%	125/125 [16:52<00:00,	Class	Images	Instances	Box(P	R
mAP50-95): 100%	18/18 [01:51					mAP50
	all	569	7670	0.855	0.66	0.745
0.481						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
17/30	0G	0.9331	0.8252	1.195	292	640:
100%	125/125 [16:32<00:00,	Class	Images	Instances	Box(P	R
mAP50-95): 100%	18/18 [01:50					mAP50
	all	569	7670	0.869	0.665	0.755
0.488						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
18/30	0G	0.9093	0.7923	1.184	275	640:
100%	125/125 [16:46<00:00,					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:50					
	all	569	7670	0.85	0.684	0.764
0.5						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
19/30	0G	0.9053	0.7639	1.17	257	640:
100%	125/125 [16:32<00:00,					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:54					
	all	569	7670	0.866	0.691	0.772
0.508						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
20/30	0G	0.8683	0.7246	1.155	349	640:
100%	125/125 [15:01<00:00,					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:47					
	all	569	7670	0.877	0.698	0.783
0.523						

Closing dataloader mosaic

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
21/30	0G	0.8941	0.6708	1.167	195	640:
100%	125/125 [14:52<00:00,					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	18/18 [01:45					
	all	569	7670	0.864	0.706	0.783
0.514						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
22/30	0G	0.869	0.6296	1.152	221	640:
100%	125/125 [14:51<00:00,					
	Class	Images	Instances	Box(P	R	mAP50
mAP50-95): 100%	18/18 [01:46					
	all	569	7670	0.875	0.718	0.796
0.538						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
23/30	0G	0.8431	0.5999	1.133	173	640:
100%	125/125 [14:50<00:00,					
	Class	Images	Instances	Box(P	R	mAP50
mAP50-95): 100%	18/18 [01:45					
	all	569	7670	0.883	0.72	0.796
0.543						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
24/30	0G	0.813	0.5655	1.108	225	640:
100%	125/125 [14:53<00:00,					
	Class	Images	Instances	Box(P	R	mAP50
mAP50-95): 100%	18/18 [01:44					
	all	569	7670	0.879	0.73	0.806
0.551						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
25/30	0G	0.7991	0.5535	1.107	213	640:
100%	125/125 [14:51<00:00,					
	Class	Images	Instances	Box(P	R	mAP50
mAP50-95): 100%	18/18 [01:43					
	all	569	7670	0.893	0.735	0.812
0.557						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	26/30	0G	0.7768	0.5278	1.09	162	640:
	125/125 [14:54<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%			18/18 [01:44				
		all	569	7670	0.891	0.737	0.815
			0.563				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	27/30	0G	0.7561	0.5144	1.078	208	640:
	125/125 [14:53<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%			18/18 [01:45				
		all	569	7670	0.892	0.74	0.817
			0.571				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	28/30	0G	0.7305	0.4915	1.058	184	640:
	125/125 [14:54<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%			18/18 [01:45				
		all	569	7670	0.897	0.75	0.819
			0.574				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	29/30	0G	0.7129	0.4776	1.052	225	640:
	125/125 [14:51<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%			18/18 [01:43				
		all	569	7670	0.9	0.748	0.823
			0.579				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	30/30	0G	0.6986	0.4668	1.047	114	640:
	125/125 [14:54<00:00,	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%			18/18 [01:43				

	all	569	7670	0.904	0.752	0.824
0.581						

30 epochs completed in 8.776 hours.

Optimizer stripped from runs\detect\new m lower batch\weights\last.pt, 52.0MB
 Optimizer stripped from runs\detect\new m lower batch\weights\best.pt, 52.0MB

Validating runs\detect\new m lower batch\weights\best.pt...

Ultralytics 8.3.32 Python-3.12.4 torch-2.5.1+cpu CPU (Intel Core(TM) i9-14900KF)

Model summary (fused): 218 layers, 25,845,550 parameters, 0 gradients, 78.7 GFLOPs

	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%		18/18 [01:29]				
0.581	all	569	7670	0.904	0.752	0.824
0.529	Hardhat	271	711	0.918	0.724	0.807
0.743	Mask	225	306	0.933	0.911	0.94
0.529	NO-Hardhat	283	454	0.93	0.702	0.806
0.399	NO-Mask	320	652	0.875	0.657	0.722
0.588	NO-Safety Vest	399	835	0.908	0.762	0.834
0.681	Person	542	2005	0.923	0.812	0.892
0.349	Safety Cone	125	618	0.835	0.563	0.656
0.601	Safety Vest	273	652	0.918	0.759	0.834
0.805	machinery	440	1114	0.94	0.9	0.945
0.591	vehicle	166	323	0.862	0.733	0.802

Speed: 0.9ms preprocess, 152.0ms inference, 0.0ms loss, 0.3ms postprocess per image

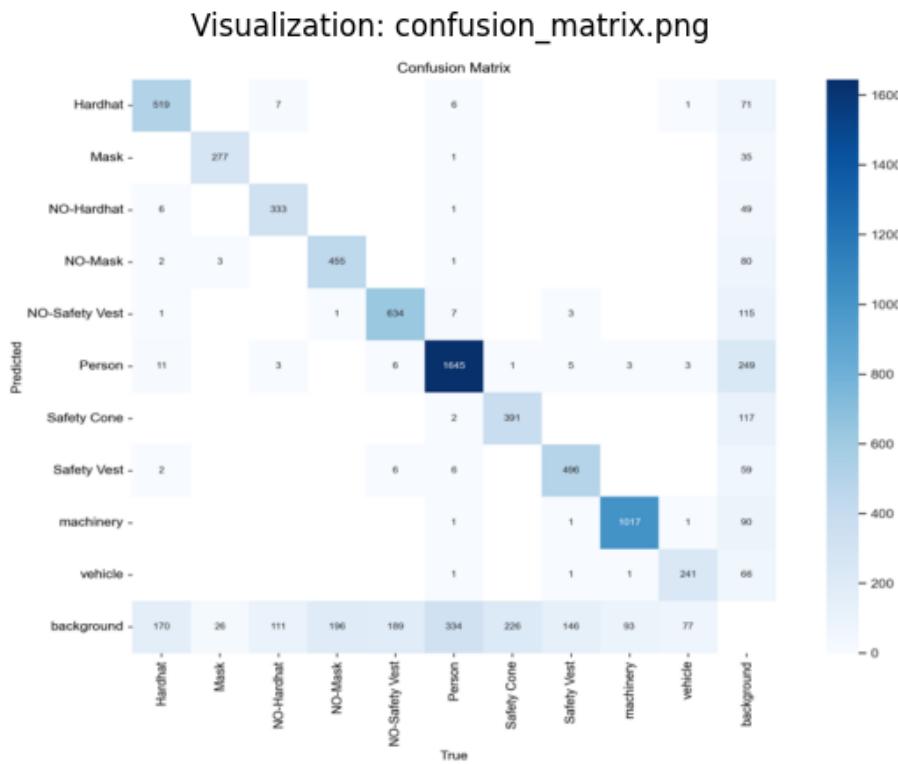
Results saved to runs\detect\new m lower batch

Total model train time: 528.16 minutes

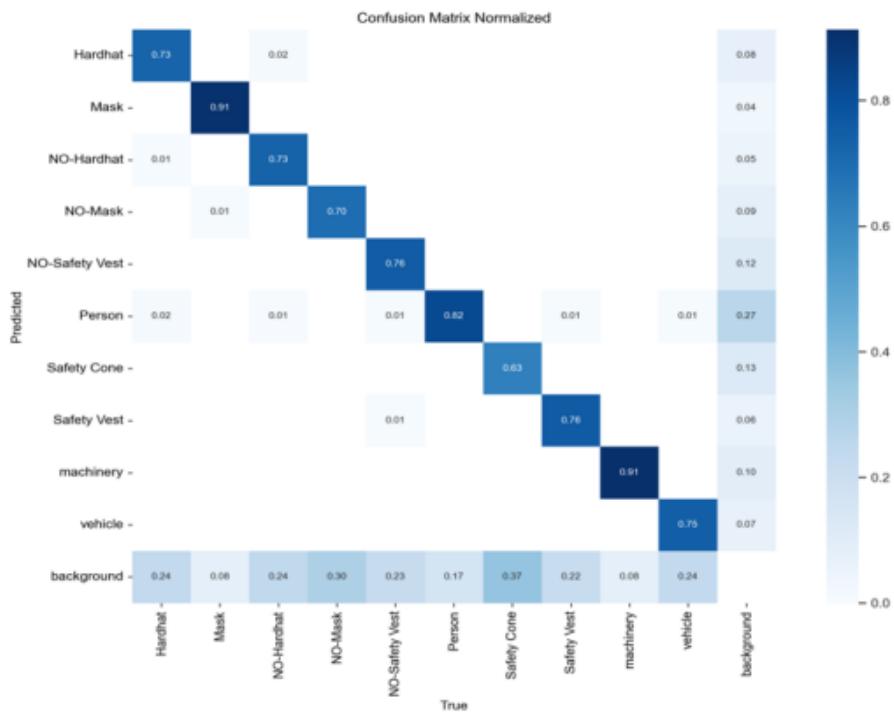
3 Model Performance Metrics

```
[7]: run_path = "runs/detect/train"
# Loop through images, do not include images for val and train
img_files = [
    f for f in os.listdir(run_path)
    if f.endswith(".png", ".jpg")) and not f.startswith(("train", "val"))
]

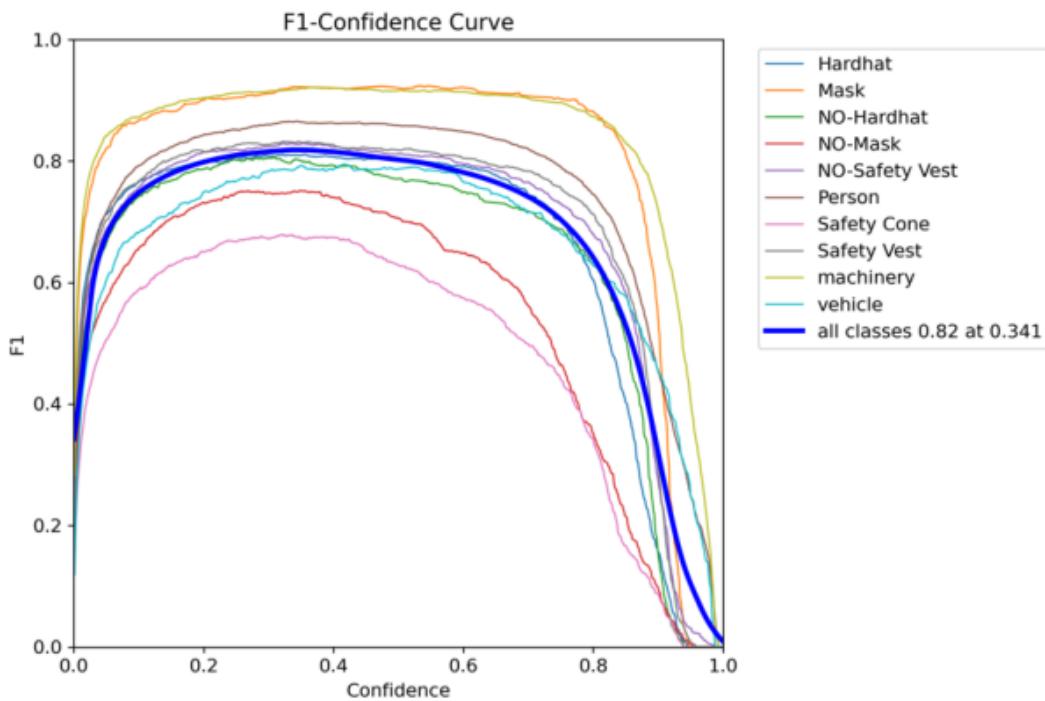
# Loop through each image file and display it
for img_file in img_files:
    file_path = os.path.join(run_path, img_file)
    image = Image.open(file_path).resize((900, 600)) # width, height
    # Display the image
    plt.figure(figsize=(8, 6))
    plt.imshow(image)
    plt.axis("off") # Hide axes
    plt.title(f"Visualization: {img_file}")
    plt.show()
```



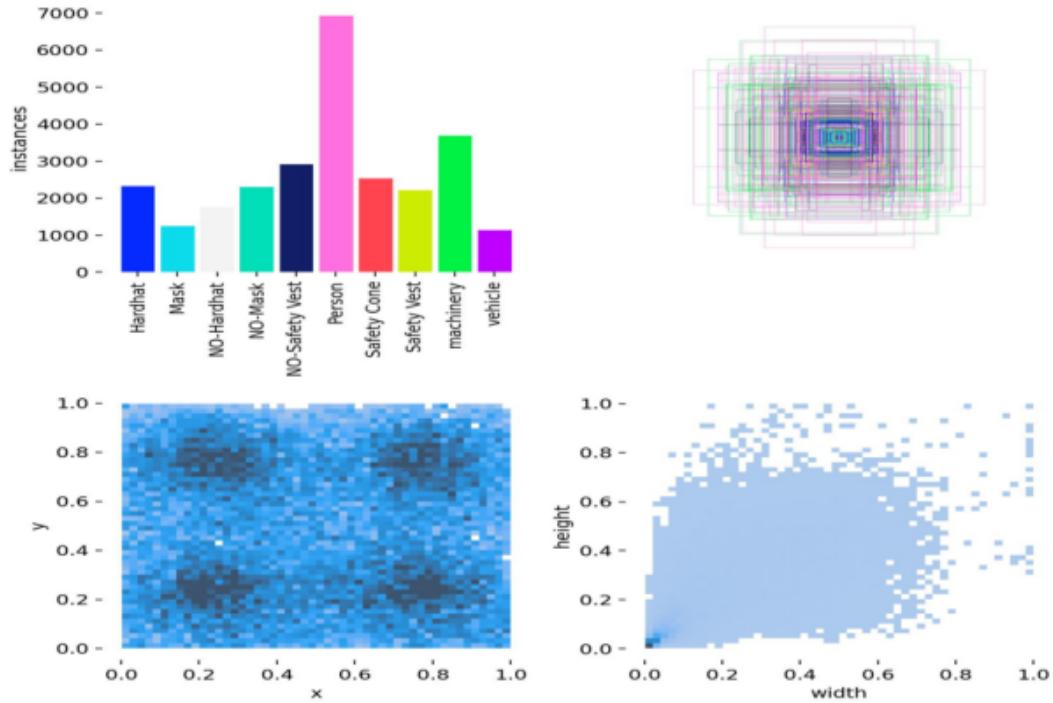
Visualization: confusion_matrix_normalized.png



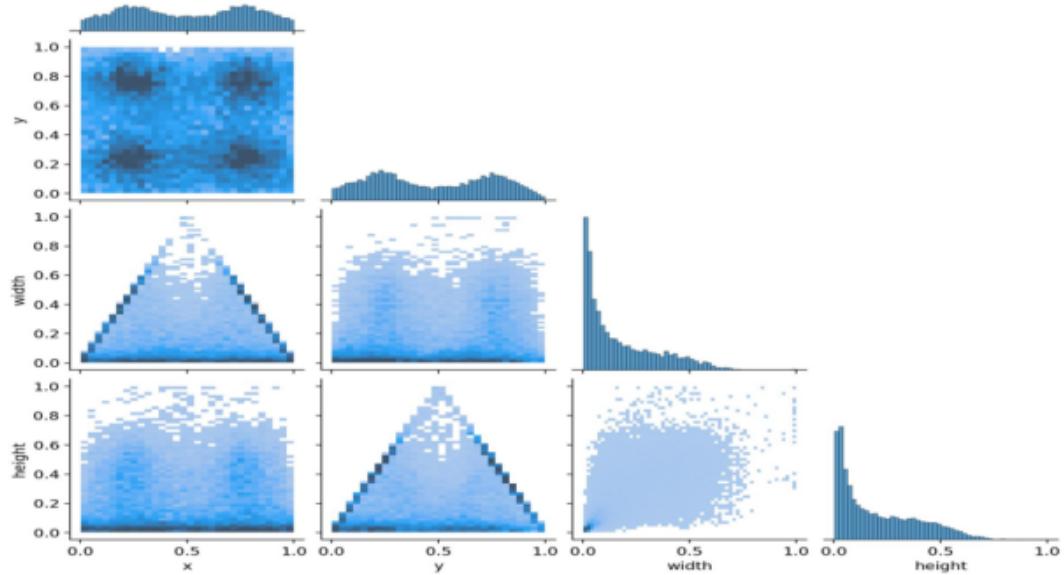
Visualization: F1_curve.png



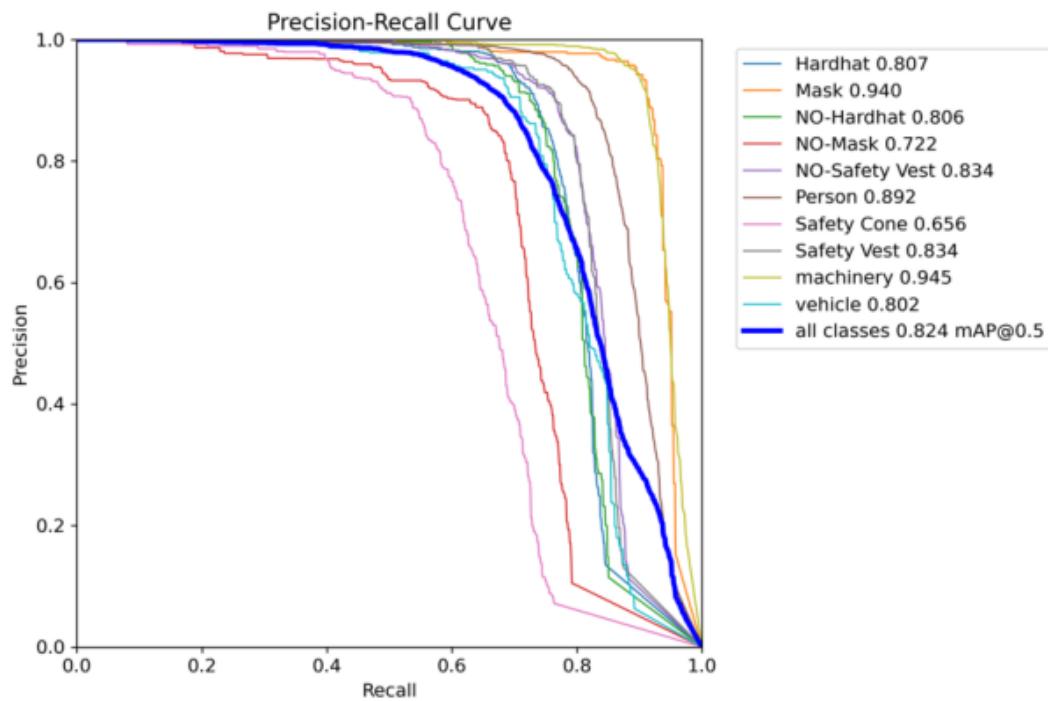
Visualization: labels.jpg



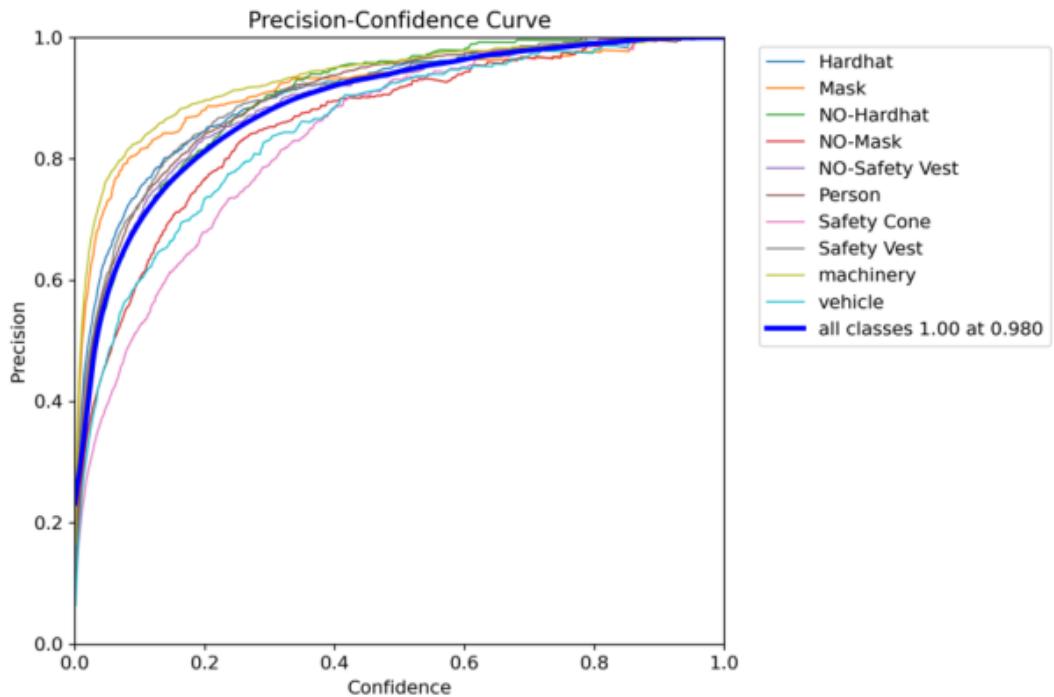
Visualization: labels_correlogram.jpg



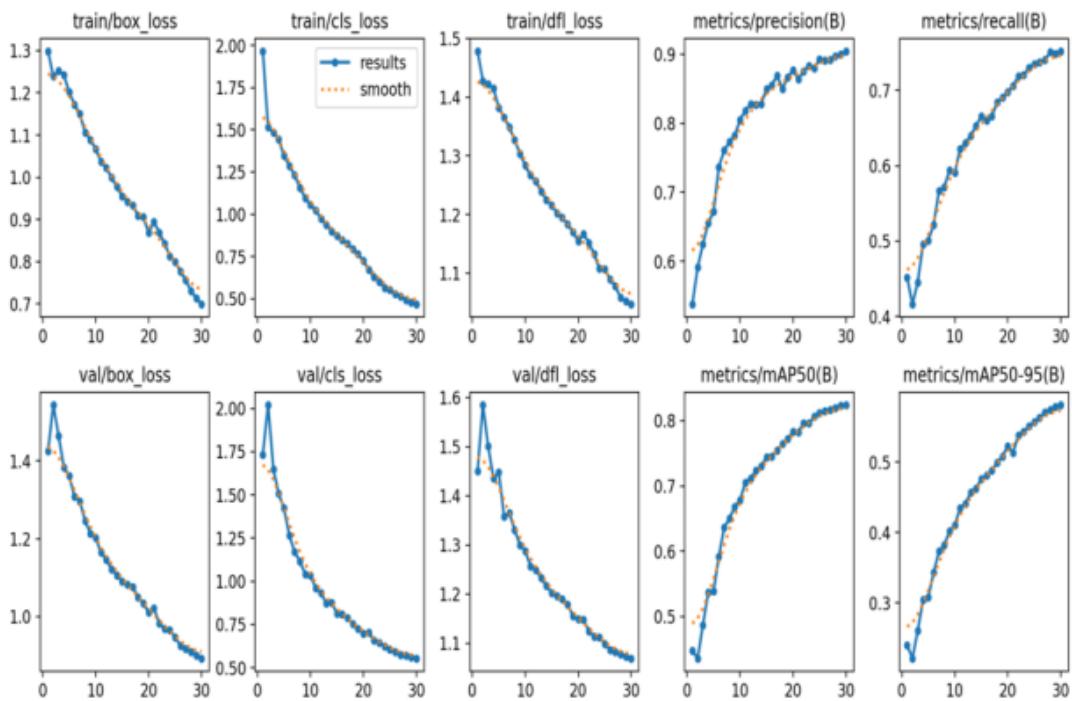
Visualization: PR_curve.png



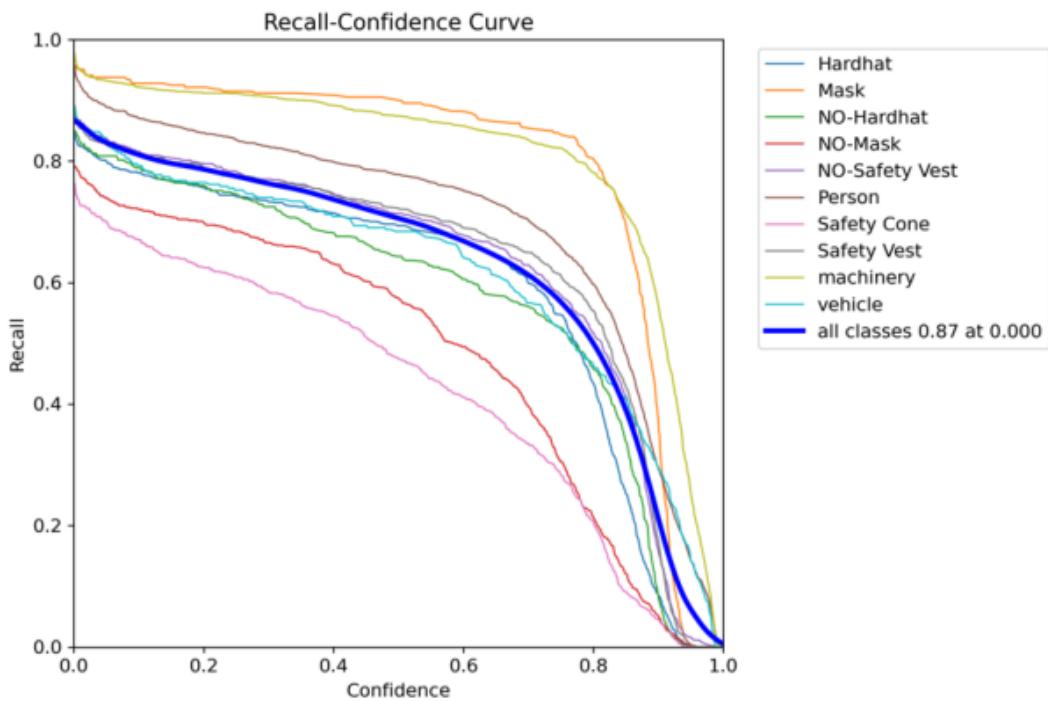
Visualization: P_curve.png



Visualization: results.png



Visualization: R_curve.png



3.0.1 Show validation image labels vs predictions

```
[28]: # Get val label and pred image
pred_image_path = glob.glob(os.path.join(run_path, "val_batch0*pred.jpg"))
labels_image_path = glob.glob(os.path.join(run_path, "val_batch0*labels.jpg"))

# Assign first matching image if found
pred_image_path = pred_image_path[0] if pred_image_path else None
labels_image_path = labels_image_path[0] if labels_image_path else None

# Error if not found
if not pred_image_path or not labels_image_path:
    print("Not enough images found!")
else:
    # Plot the images side by side
    fig, axes = plt.subplots(1, 2, figsize=(12, 6)) # 1 row, 2 columns
    for ax, image_path in zip(axes, [labels_image_path, pred_image_path]):
        image = mpimg.imread(image_path)
        ax.imshow(image)
        ax.axis('off')
        ax.set_title(os.path.basename(image_path)) # Title is the image name

    plt.tight_layout()
    plt.show()
```



3.0.2 Construction and Custom test images

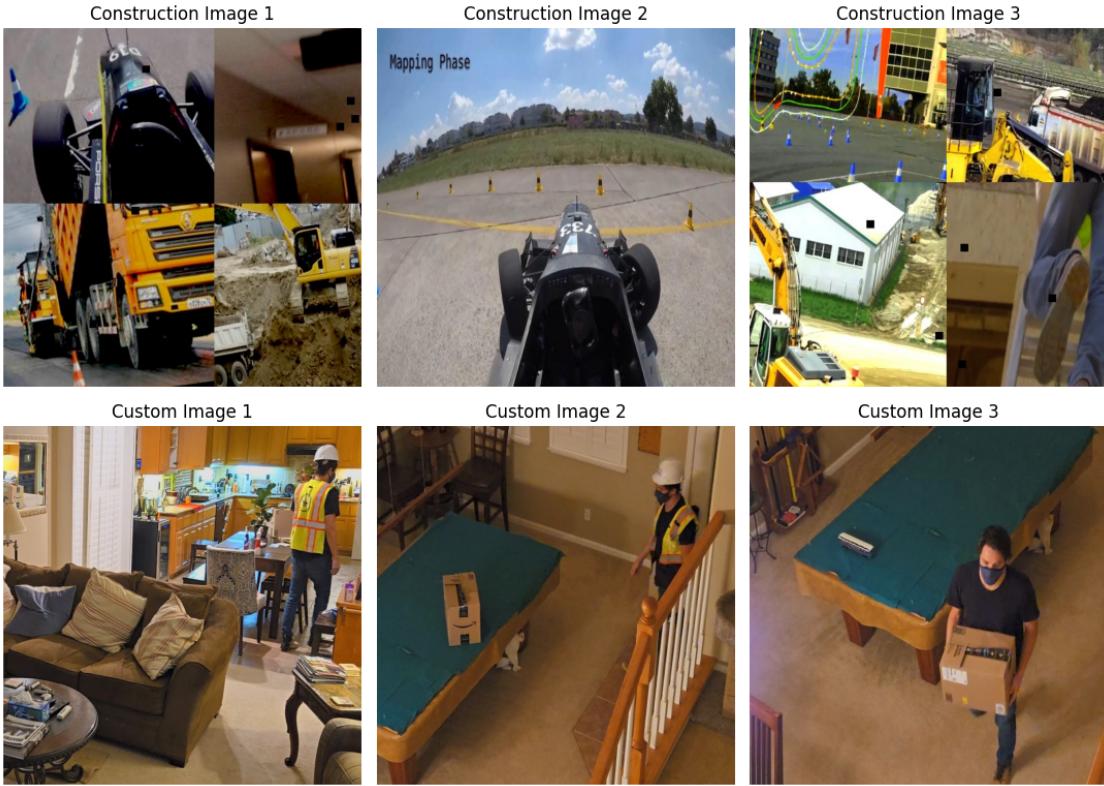
```
[36]: # Get the paths to the first three "construction" and "scene" test images
construction_images = glob.glob(f"{image_dirs['test']}/auto*.jpg")[:3]
scene_images = glob.glob(f"{image_dirs['test']}/scene*.jpg")[:3]

# Combine the image paths
image_paths = construction_images + scene_images

# Check if there are enough images
if len(image_paths) < 6:
    print("Not enough images found!")
else:
    # Create a subplot with 2 rows and 3 columns
    fig, axes = plt.subplots(2, 3, figsize=(11, 8)) # height, width

    for i, ax in enumerate(axes.flatten()):
        # Load and display each image
        image = mpimg.imread(image_paths[i])
        ax.imshow(image)
        ax.axis('off') # Hide axes
        if i < 3:
            ax.set_title(f"Construction Image {i + 1}") # Titles for construction images
        else:
            ax.set_title(f"Custom Image {i - 2}") # Titles for scene images

    plt.tight_layout()
    plt.show()
```



3.0.3 Predict test images with fine tuned model (best.pt)

```
[37]: # Remove previous predictions
detect_path = "runs/detect/"
for folder in os.listdir(detect_path):
    folder_path = os.path.join(detect_path, folder)
    if os.path.isdir(folder_path) and folder.startswith("predict"):
        shutil.rmtree(folder_path, ignore_errors=True)

# Load fine tuned model
model = YOLO("runs/detect/train/weights/best.pt")

# Make predictions on "scene" and "construction" images
for prefix in ["scene", "auto"]:
    model.predict(
        f"{image_dirs['test']}[{prefix}]*.jpg",
        save=True,           # Enable saving predictions
        conf=0.25,          # Minimum confidence threshold
        agnostic_nms=True,  # Merge overlapping boxes across classes
        iou=0.7,            # Intersection over Union threshold
        verbose=False       # Silence prediction output
    )
```

```

# Get the paths to the first three "scene" and "construction" predicted images
scene_predicted_paths = glob.glob("runs/detect/predict/scene*.jpg")[:3]
construction_predicted_paths = glob.glob("runs/detect/predict/auto*.jpg")[:3]

# Combine the image paths, construction on top row
all_predicted_paths = construction_predicted_paths + scene_predicted_paths

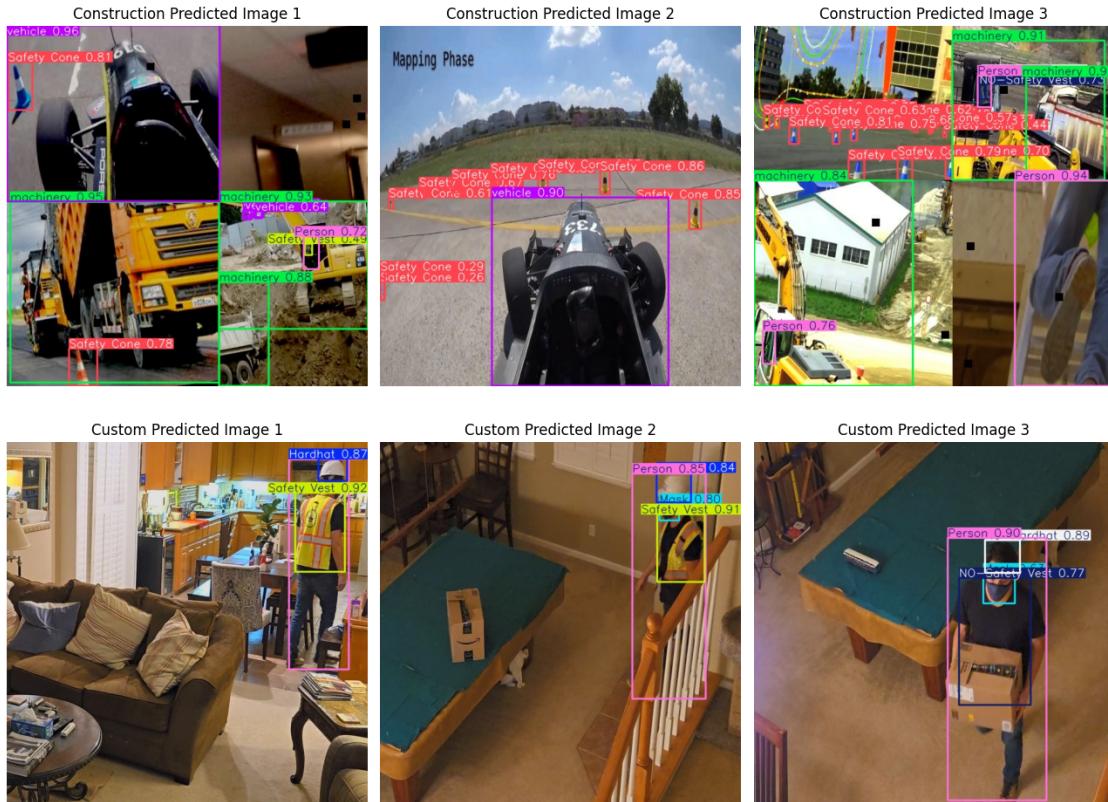
# Check if there are enough predicted images
if len(all_predicted_paths) < 6:
    print("Not enough predictions found!")
else:
    # Create a subplot with 2 rows and 3 columns
    fig, axes = plt.subplots(2, 3, figsize=(13, 10)) # Width, height

    for i, ax in enumerate(axes.flatten()):
        # Load and display each predicted image
        image = mpimg.imread(all_predicted_paths[i])
        ax.imshow(image)
        ax.axis('off') # Hide axes
        # Set titles based on the type of image
        if i < 3:
            ax.set_title(f"Construction Predicted Image {i + 1}")
        else:
            ax.set_title(f"Custom Predicted Image {i - 2}")

    plt.tight_layout()
    plt.show()

```

Results saved to runs\detect\predict
Results saved to runs\detect\predict



~/Development/aai-501-final-project/notebooks/msaai.py

```
1 # import modules
2 import os # file
3 import cv2
4 import shutil
5 import random
6 from numpy import split
7 import pandas as pd
8
9
10 # Function to get class colors
11 def getColours(cls_num):
12     base_colors = [(255, 0, 0), (0, 255, 0), (0, 0, 255)]
13     color_index = cls_num % len(base_colors)
14     increments = [(1, -2, 1), (-2, 1, -1), (1, -1, 2)]
15     color = [
16         base_colors[color_index][i]
17         + increments[color_index][i] * (cls_num // len(base_colors)) % 256
18         for i in range(3)
19     ]
20     return tuple(color)
21
22
23 def drawPredictionBox(img, class_name, box):
24     # get coordinates
25     [x1, y1, x2, y2] = box.xyxy[0]
26     # convert to int
27     x1, y1, x2, y2 = int(x1), int(y1), int(x2), int(y2)
28     vertex1 = (x1, y1)
29     vertex2 = (x2, y2)
30
31     confidence = box.conf[0]
32     colour = getColours(int(box.cls[0]))
33
34     # draw prediction rectangle
35     cv2.rectangle(img, vertex1, vertex2, colour, 2)
36
37     # put the class name and confidence on the image
38     cv2.putText(
39         img,
40         f'{class_name} {confidence:.2f}',
41         vertex1,
42         cv2.FONT_HERSHEY_SIMPLEX,
43         1,
44         colour,
45         2,
46     )
47     return img
48
49
50 def drawPredictions(predictions, img, class_names, threshold=0.4):
51     summary = ""
```

```
52     for p in predictions:
53         summary = p.verbose()
54         boxes = p.boxes
55         for box in boxes:
56             confidence = box.conf[0]
57             if confidence > threshold:
58                 cls = int(box.cls[0]) # get the class index
59                 # get the class name
60                 class_name = class_names[cls]
61                 drawPredictionBox(img, class_name, box)
62     return summary
63
64
65 # create directories to organize images and cleanup for a new to avoid
66 # duplicate images
67 def reset_directories(directories):
68     """
69     Check if the specified directories exist. If they do, delete them and
70     recreate them.
71     Ensures the directories are clean before use.
72
73     Parameters:
74         directories (list): List of directories to reset.
75     """
76     for path in directories:
77         if os.path.exists(path):
78             # delete the directory and all its contents
79             try:
80                 shutil.rmtree(path)
81                 print(f"Deleted existing directory: {path}")
82             except Exception as e:
83                 print(f"Failed to delete {path}. Reason: {e}")
84
85             # Recreate the directory
86             try:
87                 os.makedirs(path, exist_ok=True)
88                 print(f"Recreated directory: {path}")
89             except Exception as e:
90                 print(f"Failed to create directory {path}. Reason: {e}")
91
92 # SPLIT IMAGE DATASET AND FOLDERS
93 def arrange_image_and_label_files(
94     df, source_image_dir, source_label_dir, image_dirs, label_dirs
95 ):
96     for _, row in df.iterrows():
97         img_file = row["filename"]
98         split = row["split"]
99
100        # Source paths
101        img_src = os.path.join(source_image_dir, img_file)
102        label_src = os.path.join(
103            source_label_dir, os.path.splitext(img_file)[0] + ".txt"
104        )
```

```
104
105     # Destination paths
106     img_dest = os.path.join(image_dirs[split], img_file)
107     label_dest = os.path.join(
108         label_dirs[split], os.path.splitext(img_file)[0] + ".txt"
109     )
110
111     # Copy image file
112     if os.path.exists(img_src):
113         shutil.copy(img_src, img_dest)
114
115     # Copy corresponding label file
116     if os.path.exists(label_src):
117         shutil.copy(label_src, label_dest)
118
119
120 def img_train_test_split(source_image_dir):
121     """
122     Base on files in source_image_dir, split images following rule
123     Parameters:
124         source_image_dir (str): Original Images Directory to split
125     """
126     # Set random seed for reproducibility
127     random.seed(42)
128
129     # extract the image files
130     image_files = [
131         f for f in os.listdir(source_image_dir) if f.endswith((".jpg", ".jpeg",
132         ".png"))
133     ]
134     random.shuffle(image_files)
135
136     # Create DataFrame with file paths and dataset split assignments
137     split_files = pd.DataFrame({"filename": image_files})
138
139     # dynamically split dataset into train, valid, and test
140     # split sizes
141     train_size = int(0.7 * len(split_files)) # 70% for training
142     valid_size = int(0.2 * len(split_files)) # 20% for validation
143     test_size = len(split_files) - train_size - valid_size # remaining 10% for
testing
144
145     # split labels
146     train_labels = ["train"] * train_size
147     valid_labels = ["valid"] * valid_size
148     test_labels = ["test"] * test_size
149
150     split_files["split"] = train_labels + valid_labels + test_labels
151
152
153 def setup_image_directories(base_dir):
154     return
155
```

hybrid_model_train

December 8, 2024

1 Construction PPE Object Detection:

1.0.1 Hybrid Model Training

AAI-501 Group 4: Fatimat Atanda, Victor Hugo Germano, Darin Verduzco GitHub:
<https://github.com/victorhg/aai-501-final-project>

2 Dataset setup:

Download hybrid “datasets” folder into the current “hybrid_model_train” directory: Manually-annotated datasets: For images annotated manually ./datasets/manually_annotated/raw_images (178 images) ./datasets/manually_annotated/raw_labels (178 labels) Auto-annotated datasets: For images annotated with pre-trained model - Raw Images Datasets: Contains images not annotated ./datasets/raw_images (2746 images) (note: file names larger than 99 characters may not be read, update manually)

2.0.1 Import Modules and Libraries

```
[55]: import os # File
import cv2 # OpenCV for images
import yaml
import glob
import torch
import shutil
import random
import numpy as np
import pandas as pd
from PIL import Image
import seaborn as sns
from pathlib import Path
import matplotlib.image as mpimg
import matplotlib.pyplot as plt
from ultralytics import YOLO # For object detection
import albumentations as A # For image augmentation
from albumentations.pytorch import ToTensorV2 # For image formating
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_recall_curve
# Suppress all warnings
```

```
import warnings
warnings.filterwarnings("ignore")
```

2.0.2 Create directories

```
[56]: # Create direcories to organize images and cleanup for to avoid duplicates
      ↵images
dataset_base_dir = "datasets/"
# Source directory containing all raw manually annotated images
manual_source_images_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵raw_images")
manual_source_labels_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵raw_labels")
# Source directory containing all raw unannotated images
raw_images_dir = os.path.join(dataset_base_dir, "raw_images")

# Create auto-annotated directories
auto_annotated_images_dir = os.path.join(dataset_base_dir, "auto_annotated/
      ↵images")
auto_annotated_labels_dir = os.path.join(dataset_base_dir, "auto_annotated/
      ↵labels")
# Create manually annotated directories
manual_train_images_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵images/train")
manual_train_labels_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵labels/train")
manual_val_images_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵images/val")
manual_val_labels_dir = os.path.join(dataset_base_dir, "manually_annotated/
      ↵labels/val")
```

2.0.3 Reset directories

```
[57]: def reset_directories(directories):
      """
      Check if directories exist. Delete and recreate them if they do.
      Ensures the directories are clean before use.

      Parameters:
          directories (list): List of directories to reset.
      """
      for path in directories:
          if os.path.exists(path):
              # delete the directory and all its contents
              try:
                  shutil.rmtree(path)
                  print(f"Deleted existing directory: {path}")
              except Exception as e:
                  print(f"Error deleting directory: {path}. Error: {e}")
```

```

        except Exception as e:
            print(f"Failed to delete {path}. Reason: {e}")

    # Recreate the directory
    try:
        os.makedirs(path, exist_ok=True)
        print(f"Recreated directory: {path}")
    except Exception as e:
        print(f"Failed to create directory {path}. Reason: {e}")

# Define directories to reset
directories_to_reset = [
    auto_annotated_images_dir,
    auto_annotated_labels_dir,

    # manually annotated directories
    manual_train_images_dir,
    manual_train_labels_dir,
    manual_val_images_dir,
    manual_val_labels_dir
]

# Reset directories
reset_directories(directories_to_reset)
print("Directories reset and ready for use.")

# Set random seed for reproducibility
random.seed(42)

```

Deleted existing directory: datasets/auto_annotated/images
Recreated directory: datasets/auto_annotated/images
Deleted existing directory: datasets/auto_annotated/labels
Recreated directory: datasets/auto_annotated/labels
Deleted existing directory: datasets/manually_annotated/images/train
Recreated directory: datasets/manually_annotated/images/train
Deleted existing directory: datasets/manually_annotated/labels/train
Recreated directory: datasets/manually_annotated/labels/train
Deleted existing directory: datasets/manually_annotated/images/val
Recreated directory: datasets/manually_annotated/images/val
Deleted existing directory: datasets/manually_annotated/labels/val
Recreated directory: datasets/manually_annotated/labels/val
Directories reset and ready for use.

2.0.4 Create manual_data.yaml

```
[58]: # create manual_data.yaml in local datasets folder
ppe_classes = [
    'Hardhat',
    'Mask',
    'NO-Hardhat',
    'NO-Mask',
    'NO-Safety Vest',
    'Person',
    'Safety Cone',
    'Safety Vest',
    'Machinery',
    'Vehicle'
]
number_classes = len(ppe_classes)
output_dir = 'datasets'

dict_file = {
    'train': 'manually_annotated/images/train',
    'val': 'manually_annotated/images/val',
    'nc': number_classes,
    'names': ppe_classes
}

with open(os.path.join(dataset_base_dir, 'manual_data.yaml'), 'w+') as file:
    yaml.dump(dict_file, file)
# Full yaml path to prevent path error
manual_yaml_path = str(Path('datasets/manual_data.yaml').resolve())
```

2.0.5 Create final_hybrid_data.yaml

```
[59]: # create final yaml in datasets
dict_file = {
    'train': 'images/train',
    'val': 'images/val',
    'test': 'images/test',
    'nc': number_classes,
    'names': ppe_classes
}

with open(os.path.join(dataset_base_dir, 'final_hybrid_data.yaml'), 'w+') as file:
    yaml.dump(dict_file, file)
# Full yaml path to prevent path error
final_yaml_path = str(Path('datasets/final_hybrid_data.yaml').resolve())
```

2.0.6 Random Image Visualization

```
[60]: # define function to visualize random images
def visualize_images(images_dir, num_images=5):
    random.seed(42)
    image_files = [f for f in os.listdir(images_dir) if f.endswith('.jpg', '.jpeg', '.png')]
    selected_files = random.sample(image_files, min(num_images, len(image_files)))

    plt.figure(figsize=(12, 6))
    for i, image_file in enumerate(selected_files):
        image_path = os.path.join(images_dir, image_file)
        image = cv2.imread(image_path)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        plt.subplot(1, num_images, i + 1)
        plt.imshow(image)
        plt.axis("off")
        plt.title(image_file)
    plt.suptitle(f"Random Images Visualization {images_dir}")
    plt.tight_layout()
    plt.show()

visualize_images(manual_source_images_dir, num_images=2)
```



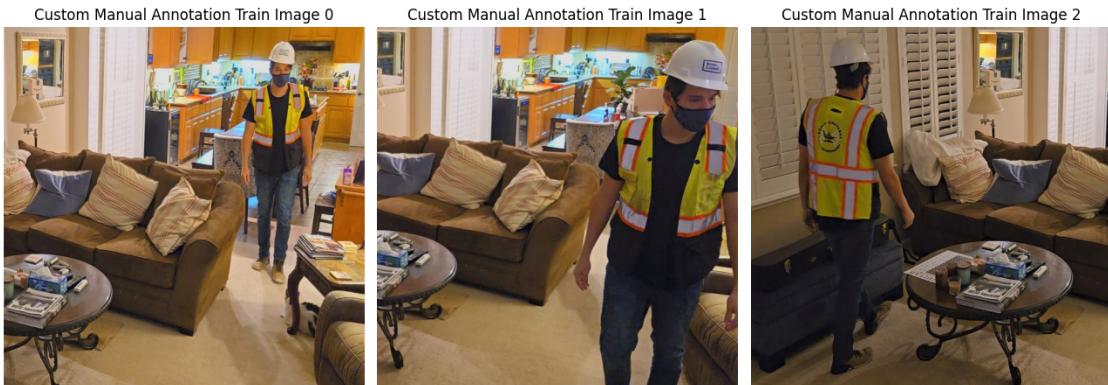
2.0.7 Custom images dataset

```
[61]: # Get the paths to the first three custom images
image_paths = glob.glob(f"{manual_source_images_dir}/scene*.jpg")[:3]

# Check if there are enough images
if len(image_paths) < 3:
    print("Not enough images found!")
else:
    # Create a subplot with 1 row and 3 columns
    fig, axes = plt.subplots(1, 3, figsize=(13, 7))  # Width, height

    for i, ax in enumerate(axes):
        # Load and display each image
        image = mpimg.imread(image_paths[i])
        ax.imshow(image)
        ax.axis('off')  # Hide axes
        ax.set_title(f"Custom Manual Annotation Train Image {i}")

    plt.tight_layout()
    plt.show()
```



2.0.8 Google images dataset

```
[62]: # Get the paths to the first three custom images
image_paths = glob.glob(f"{raw_images_dir}/images*.jpeg")[:3]

# Check if there are enough images
if len(image_paths) < 3:
    print("Not enough images found!")
else:
    # Create a subplot with 1 row and 3 columns
    fig, axes = plt.subplots(1, 3, figsize=(13, 7))  # Width, height
```

```

for i, ax in enumerate(axes):
    # Load and display each image
    image = mpimg.imread(image_paths[i])
    ax.imshow(image)
    ax.axis('off') # Hide axes
    ax.set_title(f"Custom Manual Annotation Train Image {i}")

plt.tight_layout()
plt.show()

```



2.0.9 Image Properties

```

[34]: # Define function to analyze image dimensions, aspect ratio
def analyze_image_properties(images_dir):
    properties = {"filename": [], "width": [], "height": [], "aspect_ratio": []}

    for image_file in os.listdir(images_dir):
        if image_file.endswith('.jpg', '.jpeg', '.png'):
            image_path = os.path.join(images_dir, image_file)
            image = cv2.imread(image_path)
            height, width = image.shape[:2]
            aspect_ratio = width / height
            properties["filename"].append(image_file)
            properties["width"].append(width)
            properties["height"].append(height)
            properties["aspect_ratio"].append(aspect_ratio)

    # Convert to pandas DataFrame for analysis
    df = pd.DataFrame(properties)
    return df
print("Image Properties Analysis Of Manually Labeled Images")
analyze_image_properties(manual_source_images_dir)

```

Image Properties Analysis Of Manually Labeled Images

```
[34]:
```

		filename	width	height	\
0	-1670-_png_jpg.rf.0463edb430019e01ec79eed27a63...	640	640		
1	-1670-_png_jpg.rf.3cb172ea2c4165c19ae2dd498b38...	640	640		
2	-1670-_png_jpg.rf.7da967f9aea62defc36543b9e60...	640	640		
3	-1670-_png_jpg.rf.b42b26d784545ce1a033679674a4...	640	640		
4	-1670-_png_jpg.rf.dd5cb0a4d6da02d34f1dc003fb4e...	640	640		
..		
173	youtube-619_jpg.rf.93749f0e33b2f7beb302dd2e176...	640	640		
174	youtube-697_jpg.rf.950b34f23e196850ec51a4add2e...	640	640		
175	youtube-727_jpg.rf.cb4c039b4ffea84835e67e2e3d2...	640	640		
176	youtube-804_jpg.rf.c4250b3c54e1cca8d5862297088...	640	640		
177	youtube-836_jpg.rf.df6fc9e569e7b6b850ecb75c377...	640	640		
	aspect_ratio				
0		1.0			
1		1.0			
2		1.0			
3		1.0			
4		1.0			
..		...			
173		1.0			
174		1.0			
175		1.0			
176		1.0			
177		1.0			

[178 rows x 4 columns]

2.0.10 Split manually annotated datasets to train and val

```
[35]: # Define function to split manually annotated dataset into train and validation
    ↪sets

def split_manual_dataset(images_dir, labels_dir, split_ratio=0.8):
    image_files = [
        f for f in os.listdir(images_dir) if f.lower().endswith('.jpg', '.jpeg', '.png'))
    ]
    random.shuffle(image_files)

    train_size = int(split_ratio * len(image_files))
    train_files = image_files[:train_size]
    val_files = image_files[train_size:]

    # Copy train files and corresponding labels
    for file in train_files:
        label_file = os.path.splitext(file)[0] + ".txt" # Match label with
    ↪image file
```

```

image_src = os.path.join(images_dir, file)
label_src = os.path.join(labels_dir, label_file)
shutil.copy(image_src, os.path.join(manual_train_images_dir, file))
shutil.copy(label_src, os.path.join(manual_train_labels_dir, label_file))

# Copy val files and corresponding labels
for file in val_files:
    label_file = os.path.splitext(file)[0] + ".txt" # Match label with image file
    image_src = os.path.join(images_dir, file)
    label_src = os.path.join(labels_dir, label_file)
    shutil.copy(image_src, os.path.join(manual_val_images_dir, file))
    shutil.copy(label_src, os.path.join(manual_val_labels_dir, label_file))

print("Manual dataset split completed. Images and labels copied to train and val directories.")

split_manual_dataset(manual_source_images_dir, manual_source_labels_dir)

```

Manual dataset split completed. Images and labels copied to train and val directories.

2.0.11 Train base model with manually annotated datasets

```
[ ]: # Train the generic model on manually labeled data
def train_model(manual_train_images, manual_val_images, epochs=50):

    # Train the model
    print("Training model on manually labeled dataset...")
    model = YOLO("../models/yolov8m.pt")
    model.train(
        data=manual_yaml_path,
        epochs=50, # Iterations
        imgsz=640, # Image size
        batch=16, # Batch size
        seed=88, # Train Repeatability
        lr0=1e-4, # Learning rate
        optimizer="auto", # Automatic speed, stability optimization
        patience=10 # Num of epochs to discontinue unimproved train
    )
    print("Model training completed.")
    return model

# Train the model on the manual dataset
trained_model = train_model(manual_train_images_dir, manual_val_images_dir)
```

Training model on manually labeled dataset...

```

Downloading
https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8m.pt to
'yolov8m.pt'...
100%|    | 49.7M/49.7M [00:01<00:00, 36.9MB/s]

New https://pypi.org/project/ultralytics/8.3.37 available Update with 'pip
install -U ultralytics'

Ultralytics 8.3.33 Python-3.11.9 torch-2.5.1 CPU (Apple M3 Pro)
engine/trainer: task=detect, mode=train, model=yolov8m.pt,
data=../models/manual_datasets.yaml, epochs=50, time=None, patience=10,
batch=16, imgsz=640, save=True, save_period=-1, cache=False, device=None,
workers=8, project=None, name=train, exist_ok=False, pretrained=True,
optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False,
rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0,
profile=False, freeze=None, multi_scale=False, overlap_mask=True, mask_ratio=4,
dropout=0.0, val=True, split=val, save_json=False, save_hybrid=False, conf=None,
iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None,
vid_stride=1, stream_buffer=False, visualize=False, augment=False,
agnostic_nms=False, classes=None, retina_masks=False, embed=None, show=False,
save_frames=False, save_txt=False, save_conf=False, save_crop=False,
show_labels=True, show_conf=True, show_boxes=True, line_width=None,
format=torchscript, keras=False, optimize=False, int8=False, dynamic=False,
simplify=True, opset=None, workspace=4, nms=False, lr0=0.0001, lrf=0.01,
momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8,
warmup_bias_lr=0.1, box=7.5, cls=0.5, df=1.5, pose=12.0, kobj=1.0,
label_smoothing=0.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0,
translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5,
bgr=0.0, mosaic=1.0, mixup=0.0, copy_paste=0.0, copy_paste_mode=flip,
auto_augment=randaugment, erasing=0.4, crop_fraction=1.0, cfg=None,
tracker=botsort.yaml, save_dir=runs/detect/train

Overriding model.yaml nc=80 with nc=10

```

	from	n	params	module
arguments				
0	-1	1	1392	ultralytics.nn.modules.conv.Conv
[3, 48, 3, 2]				
1	-1	1	41664	ultralytics.nn.modules.conv.Conv
[48, 96, 3, 2]				
2	-1	2	111360	ultralytics.nn.modules.block.C2f
[96, 96, 2, True]				
3	-1	1	166272	ultralytics.nn.modules.conv.Conv
[96, 192, 3, 2]				
4	-1	4	813312	ultralytics.nn.modules.block.C2f
[192, 192, 4, True]				
5	-1	1	664320	ultralytics.nn.modules.conv.Conv
[192, 384, 3, 2]				
6	-1	4	3248640	ultralytics.nn.modules.block.C2f
[384, 384, 4, True]				

```

    7           -1  1   1991808 ultralytics.nn.modules.conv.Conv
[384, 576, 3, 2]
    8           -1  2   3985920 ultralytics.nn.modules.block.C2f
[576, 576, 2, True]
    9           -1  1   831168 ultralytics.nn.modules.block.SPPF
[576, 576, 5]
    10          -1  1       0 torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']
    11          [-1, 6] 1       0 ultralytics.nn.modules.conv.Concat
[1]
    12          -1  2   1993728 ultralytics.nn.modules.block.C2f
[960, 384, 2]
    13          -1  1       0 torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']
    14          [-1, 4] 1       0 ultralytics.nn.modules.conv.Concat
[1]
    15          -1  2   517632 ultralytics.nn.modules.block.C2f
[576, 192, 2]
    16          -1  1   332160 ultralytics.nn.modules.conv.Conv
[192, 192, 3, 2]
    17          [-1, 12] 1       0 ultralytics.nn.modules.conv.Concat
[1]
    18          -1  2   1846272 ultralytics.nn.modules.block.C2f
[576, 384, 2]
    19          -1  1   1327872 ultralytics.nn.modules.conv.Conv
[384, 384, 3, 2]
    20          [-1, 9] 1       0 ultralytics.nn.modules.conv.Concat
[1]
    21          -1  2   4207104 ultralytics.nn.modules.block.C2f
[960, 576, 2]
    22          [15, 18, 21] 1   3781486 ultralytics.nn.modules.head.Detect
[10, [192, 384, 576]]
Model summary: 295 layers, 25,862,110 parameters, 25,862,094 gradients, 79.1
GFLOPs

```

Transferred 469/475 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/train',

view at <http://localhost:6006/>

Freezing layer 'model.22.dfl.conv.weight'

train: Scanning /Users/fatimatatanda/Library/CloudStorage/OneDrive-Personal/Desktop/USD/Projects/aai-501-final-project/datasets/manually_annotated/labels/train.cache... 142 images, 4 backgrounds, 0 corrupt: 100% | 142/142 [00:00<?, ?it/s]

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

```

val: Scanning /Users/fatimatatanda/Library/CloudStorage/OneDrive-
Personal/Desktop/USD/Projects/aai-501-final-
project/datasets/manually_annotated/labels/val.cache... 36 images, 0
backgrounds, 0 corrupt: 100% | 36/36 [00:00<?, ?it/s]

```

Plotting labels to runs/detect/train/labels.jpg...

```

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.0001' and
'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum'
automatically...

```

```

optimizer: AdamW(lr=0.000714, momentum=0.9) with parameter groups
77 weight(decay=0.0), 84 weight(decay=0.0005), 83 bias(decay=0.0)

```

```

TensorBoard: model graph visualization added

```

Image sizes 640 train, 640 val

Using 0 dataloader workers

Logging results to runs/detect/train

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/50	0G	1.629	3.637	1.681	297	640:
100%	9/9 [02:28<00:00, 16.49s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.83s/it]					
	all	36	402	0.68	0.167	0.194
0.135						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/50	0G	1.362	2.138	1.505	241	640:
100%	9/9 [02:27<00:00, 16.37s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.87s/it]					
	all	36	402	0.335	0.386	0.304
0.19						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
3/50	0G	1.278	1.784	1.425	193	640:
100%	9/9 [02:27<00:00, 16.35s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.75s/it]					

	all	36	402	0.567	0.384	0.4
0.248						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
4/50	0G	1.257	1.631	1.381	184	640:
100%	9/9 [02:26<00:00, 16.25s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.61s/it]					
	all	36	402	0.454	0.406	0.376
0.243						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
5/50	0G	1.218	1.443	1.351	236	640:
100%	9/9 [02:27<00:00, 16.34s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.78s/it]					
	all	36	402	0.542	0.448	0.451
0.266						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
6/50	0G	1.188	1.404	1.344	288	640:
100%	9/9 [02:28<00:00, 16.54s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.57s/it]					
	all	36	402	0.548	0.374	0.406
0.242						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
7/50	0G	1.202	1.352	1.348	188	640:
100%	9/9 [02:26<00:00, 16.31s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.60s/it]					
	all	36	402	0.423	0.48	0.414
0.241						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
8/50	0G	1.203	1.295	1.35	193	640:
100%	9/9 [02:26<00:00, 16.26s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.58s/it]					
	all	36	402	0.568	0.355	0.379
0.223						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
9/50	0G	1.244	1.372	1.387	265	640:
100%	9/9 [02:26<00:00, 16.24s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.58s/it]					
	all	36	402	0.623	0.386	0.447
0.257						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
10/50	0G	1.195	1.305	1.367	259	640:
100%	9/9 [02:27<00:00, 16.35s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.51s/it]					
	all	36	402	0.488	0.421	0.435
0.266						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
11/50	0G	1.17	1.236	1.354	220	640:
100%	9/9 [02:28<00:00, 16.55s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.59s/it]					
	all	36	402	0.582	0.407	0.454
0.262						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	12/50	0G	1.177	1.186	1.318	210	640:
	9/9	[02:27<00:00, 16.35s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:15<00:00, 7.51s/it]				mAP50
		all	36	402	0.637	0.414	0.455
							0.274

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	13/50	0G	1.138	1.163	1.304	221	640:
	9/9	[02:27<00:00, 16.42s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:14<00:00, 7.44s/it]				mAP50
		all	36	402	0.593	0.407	0.437
							0.262

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	14/50	0G	1.091	1.071	1.254	274	640:
	9/9	[02:27<00:00, 16.35s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:15<00:00, 7.57s/it]				mAP50
		all	36	402	0.785	0.38	0.458
							0.284

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	15/50	0G	1.092	1.105	1.29	232	640:
	9/9	[02:27<00:00, 16.44s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:14<00:00, 7.44s/it]				mAP50
		all	36	402	0.577	0.472	0.482
							0.269

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	16/50	0G	1.079	1.101	1.274	169	640:
	9/9	[02:27<00:00, 16.40s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:14<00:00, 7.46s/it]				mAP50

	all	36	402	0.582	0.547	0.542
0.318						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
17/50	0G	1.04	1.002	1.264	267	640:
100%	9/9 [02:26<00:00, 16.30s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	2/2 [00:14<00:00, 7.44s/it]					mAP50
	all	36	402	0.549	0.533	0.546
0.315						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
18/50	0G	1.075	1.072	1.278	198	640:
100%	9/9 [02:25<00:00, 16.19s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	2/2 [00:14<00:00, 7.42s/it]					mAP50
	all	36	402	0.703	0.407	0.493
0.291						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
19/50	0G	1.008	0.955	1.219	124	640:
100%	9/9 [02:26<00:00, 16.24s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	2/2 [00:14<00:00, 7.44s/it]					mAP50
	all	36	402	0.547	0.512	0.5
0.295						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
20/50	0G	1.008	0.9485	1.21	342	640:
100%	9/9 [02:26<00:00, 16.24s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	2/2 [00:14<00:00, 7.47s/it]					mAP50
	all	36	402	0.614	0.498	0.541
0.322						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
21/50	0G	1.024	0.9306	1.215	287	640:
100%	9/9 [13:03<00:00, 87.10s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:14<00:00, 7.34s/it]					
	all	36	402	0.705	0.528	0.578
0.336						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
22/50	0G	0.963	0.8866	1.193	284	640:
100%	9/9 [18:41<00:00, 124.64s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.70s/it]					
	all	36	402	0.745	0.488	0.576
0.36						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
23/50	0G	0.9499	0.8987	1.187	216	640:
100%	9/9 [17:55<00:00, 119.53s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.61s/it]					
	all	36	402	0.684	0.52	0.578
0.361						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
24/50	0G	0.9755	0.8833	1.183	179	640:
100%	9/9 [04:30<00:00, 30.04s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:23<00:00, 11.56s/it]					
	all	36	402	0.752	0.543	0.613
0.365						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

25/50	0G	0.9293	0.8481	1.167	224	640:
100%	9/9 [04:25<00:00, 29.47s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:23<00:00, 11.84s/it]					
	all	36	402	0.62	0.516	0.543
0.312						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
26/50	0G	0.8751	0.7917	1.148	203	640:
100%	9/9 [04:47<00:00, 31.95s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:27<00:00, 13.83s/it]					
	all	36	402	0.619	0.516	0.554
0.337						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
27/50	0G	0.8967	0.7718	1.142	334	640:
100%	9/9 [04:56<00:00, 32.99s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:25<00:00, 12.72s/it]					
	all	36	402	0.684	0.526	0.582
0.341						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
28/50	0G	0.8754	0.7799	1.131	171	640:
100%	9/9 [04:53<00:00, 32.59s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:25<00:00, 12.92s/it]					
	all	36	402	0.605	0.63	0.608
0.353						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
29/50	0G	0.86	0.7497	1.142	183	640:
100%	9/9 [04:53<00:00, 32.56s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:25<00:00, 12.93s/it]					

	all	36	402	0.718	0.594	0.639
0.385						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
30/50	0G	0.8677	0.7673	1.139	258	640:
100%	9/9 [04:47<00:00, 31.92s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:25<00:00, 12.54s/it]					
	all	36	402	0.668	0.614	0.642
0.392						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
31/50	0G	0.8359	0.7204	1.103	213	640:
100%	9/9 [04:42<00:00, 31.43s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [13:11<00:00, 395.78s/it]					
	all	36	402	0.654	0.598	0.639
0.405						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
32/50	0G	0.8033	0.6921	1.097	219	640:
100%	9/9 [05:03<00:00, 33.71s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [08:19<00:00, 249.68s/it]					
	all	36	402	0.809	0.541	0.652
0.411						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
33/50	0G	0.8198	0.6849	1.11	247	640:
100%	9/9 [02:29<00:00, 16.64s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:15<00:00, 7.75s/it]					
	all	36	402	0.649	0.641	0.655
0.424						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
34/50	0G	0.7695	0.611	1.06	194	640:
100%	9/9 [03:23<00:00, 22.56s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:17<00:00, 8.89s/it]					
	all	36	402	0.725	0.634	0.665
0.42						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
35/50	0G	0.8013	0.6653	1.103	241	640:
100%	9/9 [03:33<00:00, 23.71s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.17s/it]					
	all	36	402	0.758	0.637	0.685
0.426						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
36/50	0G	0.7606	0.6264	1.071	227	640:
100%	9/9 [03:56<00:00, 26.33s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:21<00:00, 10.99s/it]					
	all	36	402	0.826	0.609	0.685
0.423						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
37/50	0G	0.7505	0.6418	1.077	289	640:
100%	9/9 [04:00<00:00, 26.77s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.09s/it]					
	all	36	402	0.808	0.624	0.69
0.427						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

38/50	0G	0.7275	0.6043	1.048	222	640:
100%	9/9 [04:00<00:00, 26.77s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:21<00:00, 10.98s/it]					
	all	36	402	0.802	0.608	0.691
0.434						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
39/50	0G	0.7549	0.6074	1.051	171	640:
100%	9/9 [03:59<00:00, 26.62s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:21<00:00, 10.92s/it]					
	all	36	402	0.683	0.649	0.678
0.429						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
40/50	0G	0.7245	0.5878	1.053	226	640:
100%	9/9 [03:59<00:00, 26.64s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.12s/it]					
	all	36	402	0.763	0.654	0.703
0.441						

Closing dataloader mosaic
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
41/50	0G	0.739	0.6002	1.068	182	640:
100%	9/9 [04:03<00:00, 27.08s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.08s/it]					
	all	36	402	0.771	0.683	0.716
0.443						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
42/50	0G	0.7407	0.5704	1.086	169	640:
100%	9/9 [04:00<00:00, 26.70s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.15s/it]					
	all	36	402	0.796	0.668	0.724
0.444						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
43/50	0G	0.7126	0.5495	1.053	131	640:
100%	9/9 [04:06<00:00, 27.36s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.33s/it]					
	all	36	402	0.762	0.633	0.701
0.447						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
44/50	0G	0.665	0.508	1.021	137	640:
100%	9/9 [04:09<00:00, 27.76s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.28s/it]					
	all	36	402	0.754	0.628	0.697
0.45						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
45/50	0G	0.6407	0.4671	1.001	134	640:
100%	9/9 [04:12<00:00, 28.06s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	2/2 [00:22<00:00, 11.41s/it]					
	all	36	402	0.754	0.64	0.699
0.449						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	46/50	0G	0.6404	0.4704	1.005	121	640:
	9/9	[04:12<00:00, 28.07s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:22<00:00, 11.45s/it]				mAP50
		all	36	402	0.858	0.573	0.697
			0.452				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	47/50	0G	0.6175	0.448	0.9828	119	640:
	9/9	[35:16<00:00, 235.21s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:14<00:00, 7.30s/it]				mAP50
		all	36	402	0.848	0.597	0.706
			0.46				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	48/50	0G	0.6054	0.43	0.9836	144	640:
	9/9	[02:27<00:00, 16.37s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:15<00:00, 7.62s/it]				mAP50
		all	36	402	0.838	0.607	0.713
			0.469				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	49/50	0G	0.6159	0.4583	1	142	640:
	9/9	[02:40<00:00, 17.81s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:15<00:00, 7.99s/it]				mAP50
		all	36	402	0.822	0.619	0.708
			0.467				

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	50/50	0G	0.6027	0.4359	0.9766	167	640:
	9/9	[03:29<00:00, 23.25s/it]	Class	Images	Instances	Box(P)	R
mAP50-95):	100%		2/2 [00:18<00:00, 9.20s/it]				mAP50

	all	36	402	0.815	0.623	0.71
0.464						

50 epochs completed in 4.610 hours.

Optimizer stripped from runs/detect/train/weights/last.pt, 52.0MB

Optimizer stripped from runs/detect/train/weights/best.pt, 52.0MB

Validating runs/detect/train/weights/best.pt...

Ultralytics 8.3.33 Python-3.11.9 torch-2.5.1 CPU (Apple M3 Pro)

Model summary (fused): 218 layers, 25,845,550 parameters, 0 gradients, 78.7

GFLOPs

	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%		2/2 [00:18<00:00, 9.23s/it]				
0.469	all	36	402	0.839	0.607	0.712
0.525	Hardhat	21	46	0.905	0.621	0.788
0.557	Mask	17	19	0.79	0.684	0.8
0.36	NO-Hardhat	21	34	0.852	0.588	0.601
0.191	NO-Mask	26	57	0.701	0.404	0.427
0.361	NO-Safety Vest	24	53	0.869	0.5	0.647
0.579	Person	36	117	0.917	0.735	0.815
0.54	Safety Cone	7	10	0.824	0.6	0.782
0.56	Safety Vest	20	39	0.965	0.711	0.876
0.43	Machinery	11	23	0.784	0.474	0.525
0.588	Vehicle	3	4	0.784	0.75	0.856

Speed: 0.5ms preprocess, 509.7ms inference, 0.0ms loss, 0.4ms postprocess per image

Results saved to runs/detect/train

Model training completed.

2.0.12 Automatically annotate datasets with base model

```
[26]: # load the pre-trained model to auto annotated the raw_images
model_path = "runs/detect/train/weights/best.pt" # Path to your trained YOLO
↪model
def auto_annotate_images(model_path, raw_images_dir, output_images_dir, ↪
↪output_labels_dir, confidence_threshold=0.5):
    """
    Automatically annotate images using a pre-trained YOLO model.

    Args:
        - model_path (str): Path to the trained YOLO model.
        - raw_images_dir (str): Path to the directory containing raw images.
        - output_images_dir (str): Path to save annotated images.
        - output_labels_dir (str): Path to save annotation labels in YOLO format.
        - confidence_threshold (float): Minimum confidence score to include a
    ↪detection.

    """
    total_images = len([image_file for image_file in os.listdir(raw_images_dir) ↪
↪if image_file.lower().endswith('.jpg', '.jpeg', '.png')])
    # load the pre-trained model
    model = YOLO(model_path)

    # Create output directories if they don't exist
    os.makedirs(output_images_dir, exist_ok=True)
    os.makedirs(output_labels_dir, exist_ok=True)

    for i, image_file in enumerate(os.listdir(raw_images_dir)):
        if image_file.lower().endswith('.jpg', '.jpeg', '.png'):
            image_path = os.path.join(raw_images_dir, image_file)
            results = model(image_path, verbose=False) # Perform inference,
    ↪silence output for each image

            # Save the original image to the output directory
            output_image_path = os.path.join(output_images_dir, image_file)
            cv2.imwrite(output_image_path, cv2.imread(image_path))

            # Prepare the YOLO format label file
            label_file = os.path.splitext(image_file)[0] + ".txt"
            output_label_path = os.path.join(output_labels_dir, label_file)

            with open(output_label_path, "w") as f:
                for result in results[0].boxes:
                    box = result.xywhn[0].cpu().numpy() # Normalized x_center,
    ↪y_center, width, height
                    class_id = int(result.cls[0].cpu().numpy())
```

```

confidence = float(result.conf[0].cpu().numpy())

# Filter by confidence threshold
if confidence >= confidence_threshold:
    f.write(f'{class_id} {box[0]:.3f} {box[1]:.3f} {box[2]:.
˓→3f} {box[3]:.3f}\n")

# Print the current progress (overwrite previous line)
print(f"Annotating image {i+1} of {total_images}...", end='\r')

print(f"Auto-annotation completed.\n' Annotated images and labels saved in_
˓→{output_images_dir} and {output_labels_dir}.")

```

Auto-annotation completed. Annotated images and labels saved in
datasets/auto_annotated/images and datasets/auto_annotated/labels.

2.0.13 Visualize auto-annotated datasets for model accuracy

```
[39]: # visualize auto-annotation to ensure correct bounding
def visualize_annotations(image_dir, label_dir, ppe_classes, num_images=5):
    """
    Visualize YOLO-format annotations on images.

    Args:
        - image_dir (str): Path to the directory containing annotated images.
        - label_dir (str): Path to the directory containing YOLO-format label files.
        - ppe_classes (list): List of class names corresponding to class IDs.
        - num_images (int): Number of images to visualize.
    """

    num_classes = len(ppe_classes)
    colors = [
        (random.randint(0, 255),
         random.randint(0, 255),
         random.randint(0, 255)) for _ in range(num_classes)
    ]

    # Get list of image files
    image_files = [f for f in os.listdir(image_dir) if f.lower().endswith('..
˓→jpg', '.jpeg', '.png')]

    # Select a subset of images to visualize
    image_files = random.sample(image_files, min(num_images, len(image_files)))

```

```

for image_file in image_files:
    # Read the image
    image_path = os.path.join(image_dir, image_file)
    image = cv2.imread(image_path)
    if image is None:
        print(f"Could not read image: {image_file}")
        continue

    # Convert the image to RGB for visualization
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    # Read the corresponding label file
    label_file = os.path.splitext(image_file)[0] + ".txt"
    label_path = os.path.join(label_dir, label_file)
    if not os.path.exists(label_path):
        print(f"Label file not found for image: {image_file}")
        continue

    with open(label_path, "r") as f:
        labels = f.readlines()

    # Overlay bounding boxes and class names on the image
    for label in labels:
        label_data = label.strip().split()
        class_id = int(label_data[0])
        x_center, y_center, width, height = map(float, label_data[1:])

        # Convert normalized YOLO coordinates to image pixel coordinates
        img_height, img_width, _ = image.shape
        x_center *= img_width
        y_center *= img_height
        width *= img_width
        height *= img_height

        x1 = int(x_center - width / 2)
        y1 = int(y_center - height / 2)
        x2 = int(x_center + width / 2)
        y2 = int(y_center + height / 2)

        # Draw the bounding box
        #color = (255, 0, 0) # Blue color for bounding box
        color = colors[class_id]
        cv2.rectangle(image, (x1, y1), (x2, y2), color, 2)

        # Put the class name text
        label_text = f"[ppe_classes[{class_id}]}"

```

```

        cv2.putText(image, label_text, (x1, y1 - 10), cv2.
cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2)

    # Display the image
    plt.figure(figsize=(10, 8))
    plt.imshow(image)
    plt.axis("off")
    plt.title(f"Annotations for {image_file}")
    plt.show()

visualize_annotations(auto_annotated_images_dir, auto_annotated_labels_dir, n
e_classes, num_images=2)

```

2.0.14 Combine manual and auto annotated datasets

```
[40]: # Function to combine datasets (manual + auto-annotated)
# combined directory path

combined_images_dir = os.path.join(dataset_base_dir, "images")
combined_labels_dir = os.path.join(dataset_base_dir, "labels")

# check if path exist otherwise create it
os.makedirs(combined_images_dir, exist_ok=True)
os.makedirs(combined_labels_dir, exist_ok=True)

def copy_files(src_images_dir, src_labels_dir):
    for file in os.listdir(src_images_dir):
        if file.lower().endswith('.jpg', '.jpeg', '.png'):
            # Copy image file
            src_image_path = os.path.join(src_images_dir, file)
            dest_image_path = os.path.join(combined_images_dir, file)
            shutil.copy(src_image_path, dest_image_path)

            # Copy corresponding label file
            label_file = os.path.splitext(file)[0] + ".txt"
            src_label_path = os.path.join(src_labels_dir, label_file)
            dest_label_path = os.path.join(combined_labels_dir, label_file)
            shutil.copy(src_label_path, dest_label_path)

# copy manually annotated images data
print("Copying manually annotated training data...")
copy_files(manual_train_images_dir, manual_train_labels_dir)

# copy manually annotated validation data
print("Copying manually annotated validation data...")
copy_files(manual_val_images_dir, manual_val_labels_dir)
```

```

# Copy auto-annotated data
print("Copying auto-annotated data...")
copy_files(auto_annotated_images_dir, auto_annotated_labels_dir)

print("Dataset combination completed.")
print(f"Combined images stored in: {combined_images_dir}")
print(f"Combined labels stored in: {combined_labels_dir}")

```

Copying manually annotated training data...
 Copying manually annotated validation data...
 Copying auto-annotated data...
 Dataset combination completed.
 Combined images stored in: datasets/images
 Combined labels stored in: datasets/labels

2.0.15 Split combined images and labels to train, test and val

```

[41]: # extract the combined image files for random split
# split directory for detection
train_images_dir = os.path.join(dataset_base_dir, "images/train")
train_labels_dir = os.path.join(dataset_base_dir, "labels/train")

test_images_dir = os.path.join(dataset_base_dir, "images/test")
test_labels_dir = os.path.join(dataset_base_dir, "labels/test")

val_images_dir = os.path.join(dataset_base_dir, "images/val")
val_labels_dir = os.path.join(dataset_base_dir, "labels/val")

# create dataset directories for detection & prediction
for dir_path in [
    train_images_dir,
    train_labels_dir,
    test_images_dir,
    test_labels_dir,
    val_images_dir,
    val_labels_dir
]:
    os.makedirs(dir_path, exist_ok=True)

image_files = [f for f in os.listdir(combined_images_dir) if f.endswith('.jpg', '.jpeg', '.png')]

random.shuffle(image_files)

# Create DataFrame with file paths and dataset split assignments
df = pd.DataFrame({"filename": image_files})
# dynamically split dataset into train, valid, and test

```

```

# split sizes
train_size = int(0.7 * len(df)) # 70% for training
val_size = int(0.2 * len(df)) # 20% for validation
test_size = len(df) - train_size - val_size # remaining 10% for testing

# split labels
train_labels = ["train"] * train_size
val_labels = ["val"] * val_size
test_labels = ["test"] * test_size

# combine and assign to DataFrame
df["split"] = train_labels + val_labels + test_labels

# Function to move files into corresponding split directories
def move_files(df, src_images_dir, src_labels_dir, dest_images_dir,dest_labels_dir):
    for _, row in df.iterrows(): # iterate over a DataFrame row by row
        filename = row["filename"]
        split = row["split"]
        image_path = os.path.join(src_images_dir, filename)
        label_path = os.path.join(src_labels_dir, os.path.splitext(filename)[0] + ".txt")

        dest_image_dir = dest_images_dir[split]
        dest_label_dir = dest_labels_dir[split]

        # move image
        shutil.move(image_path, os.path.join(dest_image_dir, filename))

        # move label
        shutil.move(label_path, os.path.join(dest_label_dir, os.path.splitext(filename)[0] + ".txt"))

# map split directories for images and labels
dest_images_dir = {"train": train_images_dir, "val": val_images_dir, "test":test_images_dir}
dest_labels_dir = {"train": train_labels_dir, "val": val_labels_dir, "test":test_labels_dir}

# call move fuxn
move_files(df, combined_images_dir, combined_labels_dir, dest_images_dir,dest_labels_dir)

print("Dataset split completed.")

```

```

print(f"Training images: {len(os.listdir(train_images_dir))}, Labels: {len(os.
    ↪listdir(train_labels_dir))}")
print(f"Validation images: {len(os.listdir(val_images_dir))}, Labels: {len(os.
    ↪listdir(val_labels_dir))}")
print(f"Testing images: {len(os.listdir(test_images_dir))}, Labels: {len(os.
    ↪listdir(test_labels_dir))}")

```

Dataset split completed.

Training images: 2043, Labels: 2043

Validation images: 584, Labels: 584

Testing images: 293, Labels: 293

2.0.16 Train the final model with the combined datasets

(NOTE: update model_path for the base model created)

```

[33]: input_name = input('Please enter run folder description:')
# load pre-trained YOLO base model
model_path = "runs/detect/train/weights/best.pt"
model = YOLO(model_path)
# Train the model
print("Training the final model...")
model.train(
    data=final_yaml_path,
    epochs=30, # Iterations
    imgsz=640, # Image size
    batch=20, # Batch size
    lr0=1e-3, # Learning rate
    patience=10, # Num of epochs to discontinue unimproved train
    amp=True, # Reduce memore, increase speed
    plots = True, # produce plots
    seed=88, # Train repeatability
    name = input_name, # Run folder name
)
print("Model training completed.")

```

Training the final model...

New <https://pypi.org/project/ultralytics/8.3.38> available Update with 'pip install -U ultralytics'

Ultralytics 8.3.33 Python-3.11.9 torch-2.5.1 CPU (Apple M3 Pro)

```

engine/trainer: task=detect, mode=train,
model=runs/detect/train/weights/best.pt, data=../models/datasets.yaml,
epochs=30, time=None, patience=10, batch=20, imgsz=640, save=True,
save_period=-1, cache=False, device=None, workers=8, project=None, name=train4,
exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0,
deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10,
resume=False, amp=True, fraction=1.0, profile=False, freeze=None,
multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True,
split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300,

```

```

half=False, dnn=False, plots=True, source=None, vid_stride=1,
stream_buffer=False, visualize=False, augment=False, agnostic_nms=False,
classes=None, retina_masks=False, embed=None, show=False, save_frames=False,
save_txt=False, save_conf=False, save_crop=False, show_labels=True,
show_conf=True, show_boxes=True, line_width=None, format=torchscript,
keras=False, optimize=False, int8=False, dynamic=False, simplify=True,
opset=None, workspace=4, nms=False, lr0=0.001, lrf=0.01, momentum=0.937,
weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1,
box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, label_smoothing=0.0, nbs=64,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5,
shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0,
mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment,
erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml,
save_dir=runs/detect/train4

```

	from	n	params	module
arguments				
0		-1 1	1392	ultralytics.nn.modules.conv.Conv
[3, 48, 3, 2]				
1		-1 1	41664	ultralytics.nn.modules.conv.Conv
[48, 96, 3, 2]				
2		-1 2	111360	ultralytics.nn.modules.block.C2f
[96, 96, 2, True]				
3		-1 1	166272	ultralytics.nn.modules.conv.Conv
[96, 192, 3, 2]				
4		-1 4	813312	ultralytics.nn.modules.block.C2f
[192, 192, 4, True]				
5		-1 1	664320	ultralytics.nn.modules.conv.Conv
[192, 384, 3, 2]				
6		-1 4	3248640	ultralytics.nn.modules.block.C2f
[384, 384, 4, True]				
7		-1 1	1991808	ultralytics.nn.modules.conv.Conv
[384, 576, 3, 2]				
8		-1 2	3985920	ultralytics.nn.modules.block.C2f
[576, 576, 2, True]				
9		-1 1	831168	ultralytics.nn.modules.block.SPPF
[576, 576, 5]				
10		-1 1	0	torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']				
11		[-1, 6] 1	0	ultralytics.nn.modules.conv.Concat
[1]				
12		-1 2	1993728	ultralytics.nn.modules.block.C2f
[960, 384, 2]				
13		-1 1	0	torch.nn.modules.upsampling.Upsample
[None, 2, 'nearest']				
14		[-1, 4] 1	0	ultralytics.nn.modules.conv.Concat
[1]				
15		-1 2	517632	ultralytics.nn.modules.block.C2f

```

[576, 192, 2]
16           -1  1    332160 ultralytics.nn.modules.conv.Conv
[192, 192, 3, 2]
17      [-1, 12]  1        0 ultralytics.nn.modules.conv.Concat
[1]
18           -1  2    1846272 ultralytics.nn.modules.block.C2f
[576, 384, 2]
19           -1  1    1327872 ultralytics.nn.modules.conv.Conv
[384, 384, 3, 2]
20      [-1, 9]  1        0 ultralytics.nn.modules.conv.Concat
[1]
21           -1  2    4207104 ultralytics.nn.modules.block.C2f
[960, 576, 2]
22      [15, 18, 21]  1    3781486 ultralytics.nn.modules.head.Detect
[10, [192, 384, 576]]
Model summary: 295 layers, 25,862,110 parameters, 25,862,094 gradients, 79.1
GFLOPs

```

Transferred 475/475 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/train4',

view at <http://localhost:6006/>

Freezing layer 'model.22.dfl.conv.weight'

train: Scanning /Users/fatimatatanda/Library/CloudStorage/OneDrive-Personal/Desktop/USD/Projects/aai-501-final-project/datasets/labels/train.cache... 2656 images, 60 backgrounds, 0 corrupt: 100% | 2656/2656 [00:00<?, ?it/s]

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

val: Scanning /Users/fatimatatanda/Library/CloudStorage/OneDrive-Personal/Desktop/USD/Projects/aai-501-final-project/datasets/labels/val.cache... 175 images, 1 backgrounds, 0 corrupt: 100% | 175/175 [00:00<?, ?it/s]

Plotting labels to runs/detect/train4/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.001' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

optimizer: AdamW(lr=0.000714, momentum=0.9) with parameter groups
77 weight(decay=0.0), 84 weight(decay=0.00046875), 83 bias(decay=0.0)

TensorBoard: model graph visualization added

Image sizes 640 train, 640 val

Using 0 dataloader workers

Logging results to runs/detect/train4

Starting training for 30 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/30	0G	0.8476	1.012	1.197	221	640:
100%	133/133 [1:17:13<00:00, 34.84s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.88s/it]					
	all	175	1409	0.751	0.762	0.825
0.62						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/30	0G	0.9253	1.116	1.239	243	640:
100%	133/133 [45:30<00:00, 20.53s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.86s/it]					
	all	175	1409	0.693	0.655	0.692
0.462						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
3/30	0G	0.9616	1.176	1.26	143	640:
100%	133/133 [45:34<00:00, 20.56s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.82s/it]					
	all	175	1409	0.631	0.664	0.675
0.444						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
4/30	0G	0.9604	1.157	1.255	200	640:
100%	133/133 [45:32<00:00, 20.55s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.85s/it]					
	all	175	1409	0.714	0.705	0.759
0.516						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

	5/30	0G	0.9396	1.12	1.242	224	640:
100%	133/133 [45:31<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.85s/it]	all	175	1409	0.728	0.733	0.769
0.525							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
6/30	0G	0.9102	1.081	1.222	156	640:	
100%	133/133 [45:32<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.81s/it]	all	175	1409	0.722	0.759	0.806
0.576							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
7/30	0G	0.8813	1.02	1.205	196	640:	
100%	133/133 [45:34<00:00, 20.56s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.88s/it]	all	175	1409	0.759	0.745	0.797
0.562							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
8/30	0G	0.8526	0.9793	1.189	155	640:	
100%	133/133 [45:31<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.87s/it]	all	175	1409	0.771	0.733	0.79
0.56							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
9/30	0G	0.8388	0.9419	1.183	200	640:	
100%	133/133 [45:30<00:00, 20.53s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.84s/it]	all	175	1409	0.771	0.733	0.79

	all	175	1409	0.747	0.773	0.815
0.601						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
10/30	0G	0.8247	0.9245	1.17	225	640:
100%	133/133 [45:30<00:00, 20.53s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5 [01:09<00:00, 13.86s/it]					mAP50
	all	175	1409	0.756	0.786	0.819
0.597						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
11/30	0G	0.8024	0.8924	1.16	198	640:
100%	133/133 [45:32<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5 [01:09<00:00, 13.88s/it]					mAP50
	all	175	1409	0.8	0.787	0.847
0.623						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
12/30	0G	0.794	0.8722	1.146	180	640:
100%	133/133 [45:33<00:00, 20.55s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5 [01:08<00:00, 13.79s/it]					mAP50
	all	175	1409	0.77	0.79	0.834
0.609						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
13/30	0G	0.7835	0.848	1.139	303	640:
100%	133/133 [45:33<00:00, 20.55s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5 [01:09<00:00, 13.84s/it]					mAP50
	all	175	1409	0.817	0.788	0.852
0.639						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
14/30	0G	0.7612	0.8223	1.127	298	640:
100%	133/133 [45:31<00:00, 20.54s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.85s/it]					
	all	175	1409	0.804	0.789	0.858
0.645						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
15/30	0G	0.7546	0.806	1.122	160	640:
100%	133/133 [45:32<00:00, 20.55s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.86s/it]					
	all	175	1409	0.791	0.803	0.851
0.65						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
16/30	0G	0.7356	0.7834	1.114	171	640:
100%	133/133 [45:33<00:00, 20.55s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:08<00:00, 13.79s/it]					
	all	175	1409	0.811	0.819	0.866
0.656						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
17/30	0G	0.7242	0.7642	1.1	252	640:
100%	133/133 [45:32<00:00, 20.55s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:08<00:00, 13.78s/it]					
	all	175	1409	0.805	0.837	0.868
0.662						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	18/30	0G	0.7128	0.7491	1.101	305	640:
	133/133 [45:31<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.88s/it]	all	175	1409	0.813	0.822	0.867
	0.666						

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	19/30	0G	0.7043	0.7331	1.094	190	640:
	133/133 [45:31<00:00, 20.54s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.85s/it]	all	175	1409	0.821	0.818	0.868
	0.672						

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	20/30	0G	0.6807	0.7081	1.079	157	640:
	133/133 [45:32<00:00, 20.55s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.82s/it]	all	175	1409	0.802	0.845	0.883
	0.685						

Closing dataloader mosaic
albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	21/30	0G	0.6734	0.6241	1.082	100	640:
	133/133 [45:28<00:00, 20.52s/it]	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.84s/it]	all	175	1409	0.814	0.834	0.868
	0.674						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
22/30	0G	0.6487	0.5883	1.067	99	640:
100%	133/133 [45:28<00:00, 20.51s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.87s/it]					
	all	175	1409	0.798	0.861	0.886
0.694						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
23/30	0G	0.6351	0.5559	1.058	119	640:
100%	133/133 [45:36<00:00, 20.58s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.81s/it]					
	all	175	1409	0.816	0.823	0.883
0.687						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
24/30	0G	0.6246	0.5412	1.052	103	640:
100%	133/133 [45:31<00:00, 20.54s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.87s/it]					
	all	175	1409	0.816	0.826	0.882
0.702						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
25/30	0G	0.601	0.5232	1.035	117	640:
100%	133/133 [47:50<00:00, 21.58s/it]					
	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%	5/5 [01:09<00:00, 13.82s/it]					
	all	175	1409	0.824	0.821	0.876
0.693						

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	26/30	0G	0.5886	0.5004	1.033	85	640:
	133/133	[45:27<00:00, 20.51s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5	[01:08<00:00, 13.73s/it]					mAP50
	all		175	1409	0.825	0.85	0.884
0.704							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	27/30	0G	0.5779	0.4901	1.025	167	640:
	133/133	[45:27<00:00, 20.51s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5	[01:09<00:00, 13.86s/it]					mAP50
	all		175	1409	0.801	0.85	0.887
0.704							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	28/30	0G	0.5602	0.4692	1.011	96	640:
	133/133	[45:36<00:00, 20.57s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5	[01:09<00:00, 13.84s/it]					mAP50
	all		175	1409	0.83	0.832	0.883
0.706							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	29/30	0G	0.5486	0.4525	1.005	77	640:
	133/133	[45:37<00:00, 20.58s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5	[01:09<00:00, 13.90s/it]					mAP50
	all		175	1409	0.804	0.868	0.89
0.715							

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100%	30/30	0G	0.5328	0.435	0.9977	136	640:
	133/133	[45:38<00:00, 20.59s/it]	Class	Images	Instances	Box(P)	R
mAP50-95): 100%	5/5	[01:09<00:00, 13.88s/it]					mAP50

	all	175	1409	0.844	0.843	0.89
0.714						

30 epochs completed in 23.917 hours.

Optimizer stripped from runs/detect/train4/weights/last.pt, 52.0MB

Optimizer stripped from runs/detect/train4/weights/best.pt, 52.0MB

Validating runs/detect/train4/weights/best.pt...

Ultralytics 8.3.33 Python-3.11.9 torch-2.5.1 CPU (Apple M3 Pro)

Model summary (fused): 218 layers, 25,845,550 parameters, 0 gradients, 78.7

GFLOPs

	Class	Images	Instances	Box(P)	R	mAP50
mAP50-95): 100%		5/5 [01:06<00:00, 13.37s/it]				
0.715	all	175	1409	0.805	0.868	0.89
0.83	Hardhat	50	114	0.863	0.956	0.972
0.782	Mask	60	74	0.829	0.946	0.929
0.737	NO-Hardhat	63	88	0.833	0.875	0.915
0.56	NO-Mask	55	100	0.772	0.83	0.841
0.62	NO-Safety Vest	80	128	0.722	0.73	0.808
0.825	Person	150	412	0.885	0.92	0.956
0.744	Safety Cone	29	99	0.89	0.96	0.949
0.709	Safety Vest	61	114	0.785	0.842	0.89
0.739	Machinery	123	247	0.778	0.822	0.854
0.601	Vehicle	23	33	0.688	0.801	0.783

Speed: 0.5ms preprocess, 376.9ms inference, 0.0ms loss, 0.2ms postprocess per image

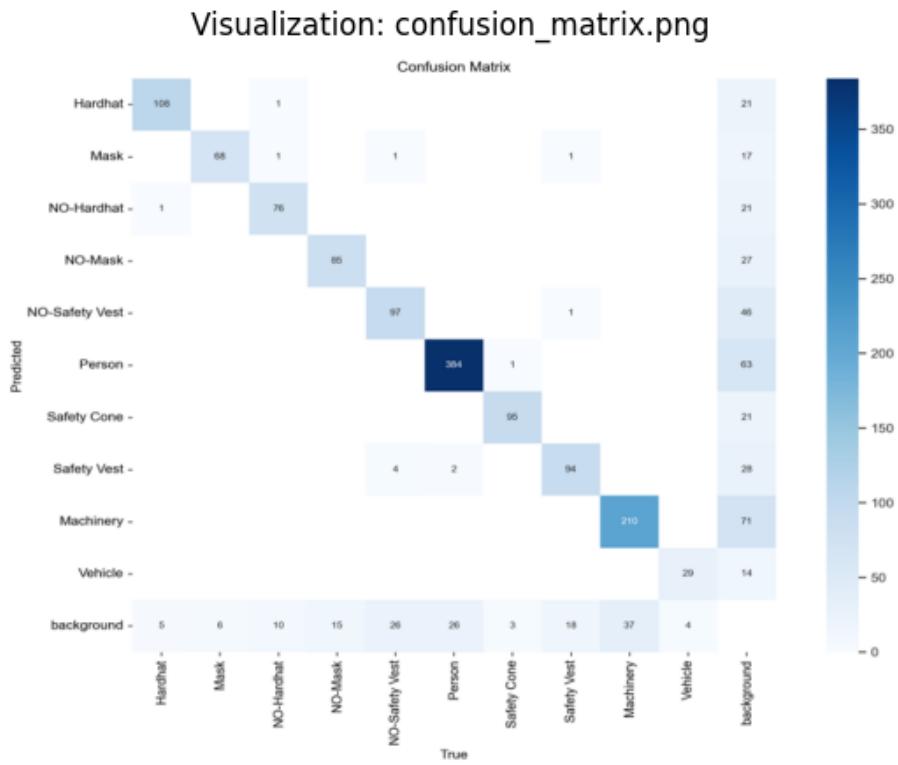
Results saved to runs/detect/train4

Model training completed.

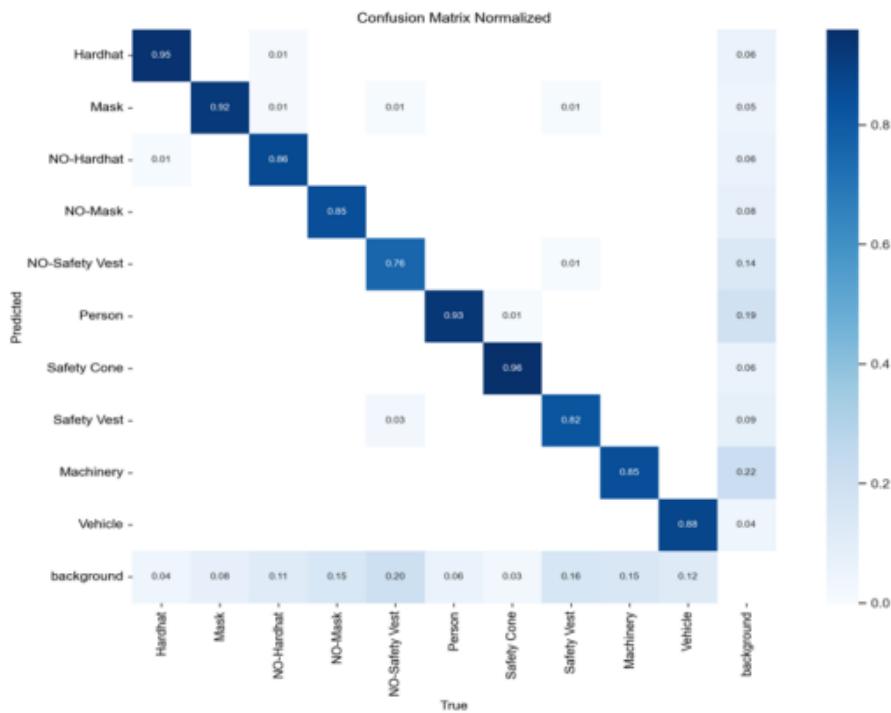
2.0.17 Evaluate train results

```
[52]: run_path = "runs/detect/train4"
# loop through images, do not include images for val and train
img_files = [
    f for f in os.listdir(run_path)
    if f.endswith(".png", ".jpg")) and not f.startswith(("train", "val"))
]

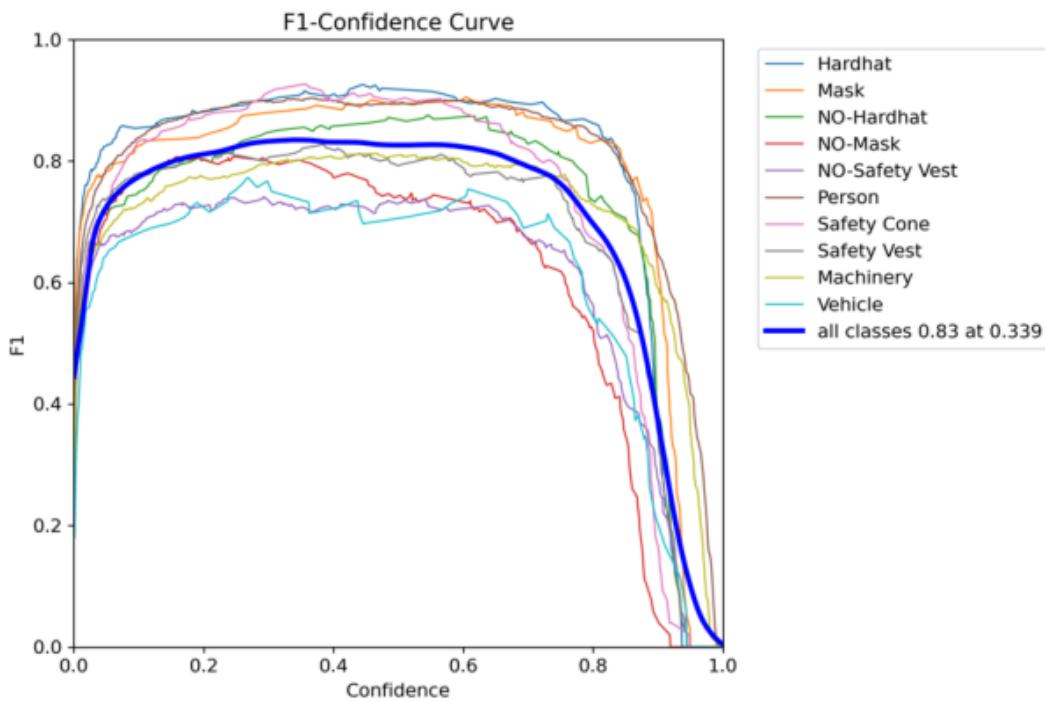
# Loop through each image file and display it
for img_file in img_files:
    file_path = os.path.join(run_path, img_file)
    image = Image.open(file_path).resize((900, 600)) # Width, height
    # Display the image
    plt.figure(figsize=(8, 6))
    plt.imshow(image)
    plt.axis("off") # Hide axes
    plt.title(f"Visualization: {img_file}")
    plt.show()
```



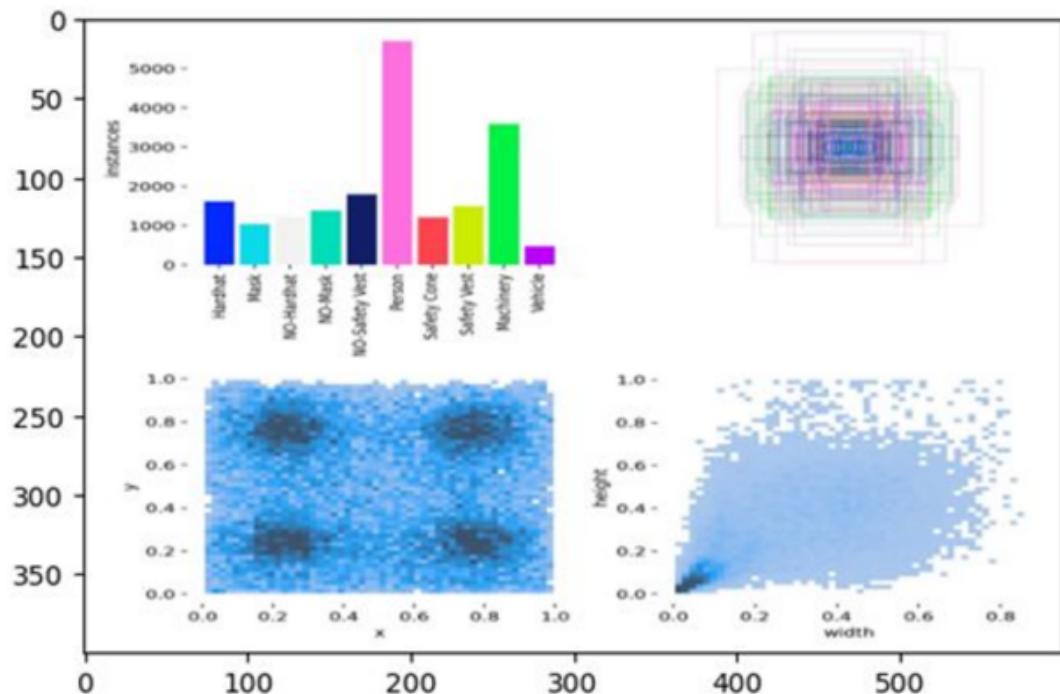
Visualization: confusion_matrix_normalized.png



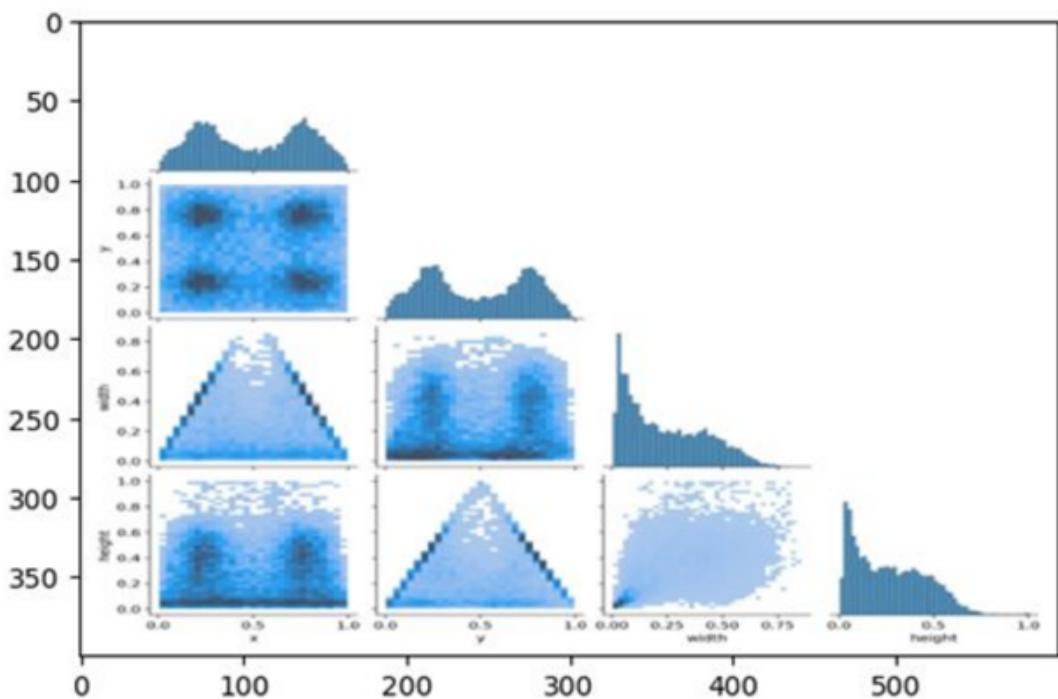
Visualization: F1_curve.png



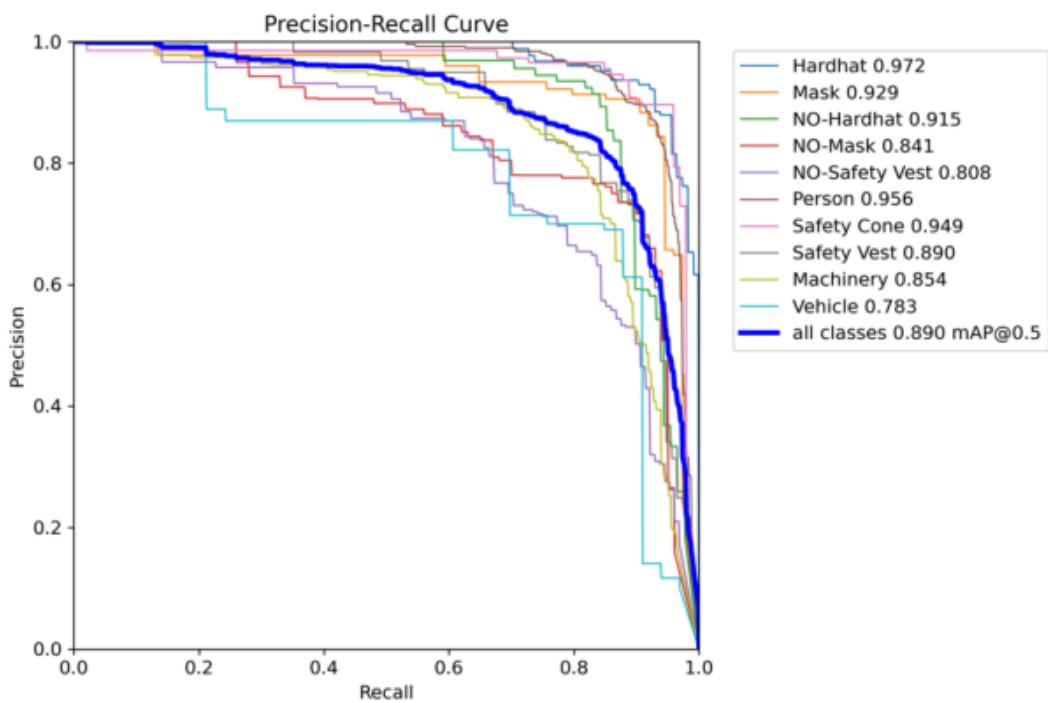
Visualization: labels.jpg



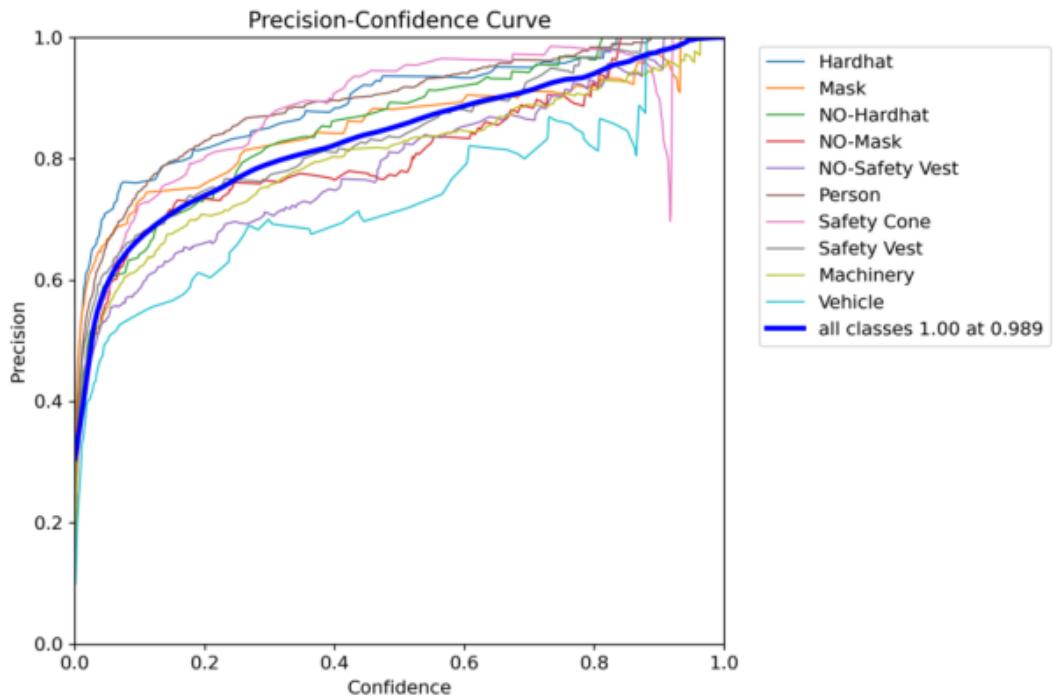
Visualization: labels_correlogram.jpg



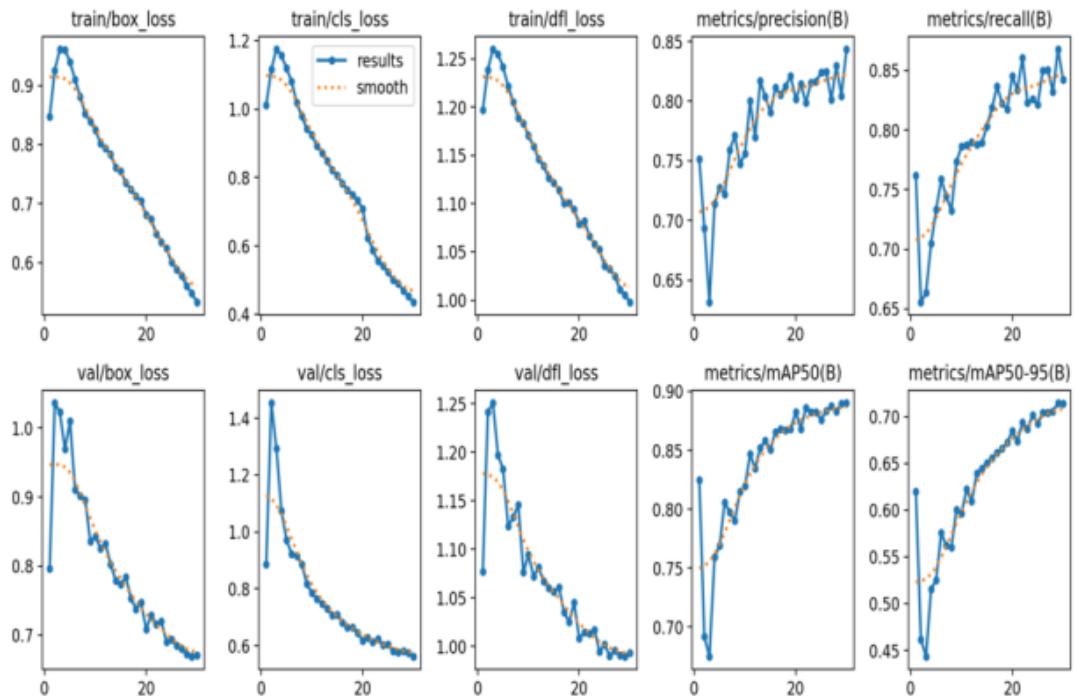
Visualization: PR_curve.png



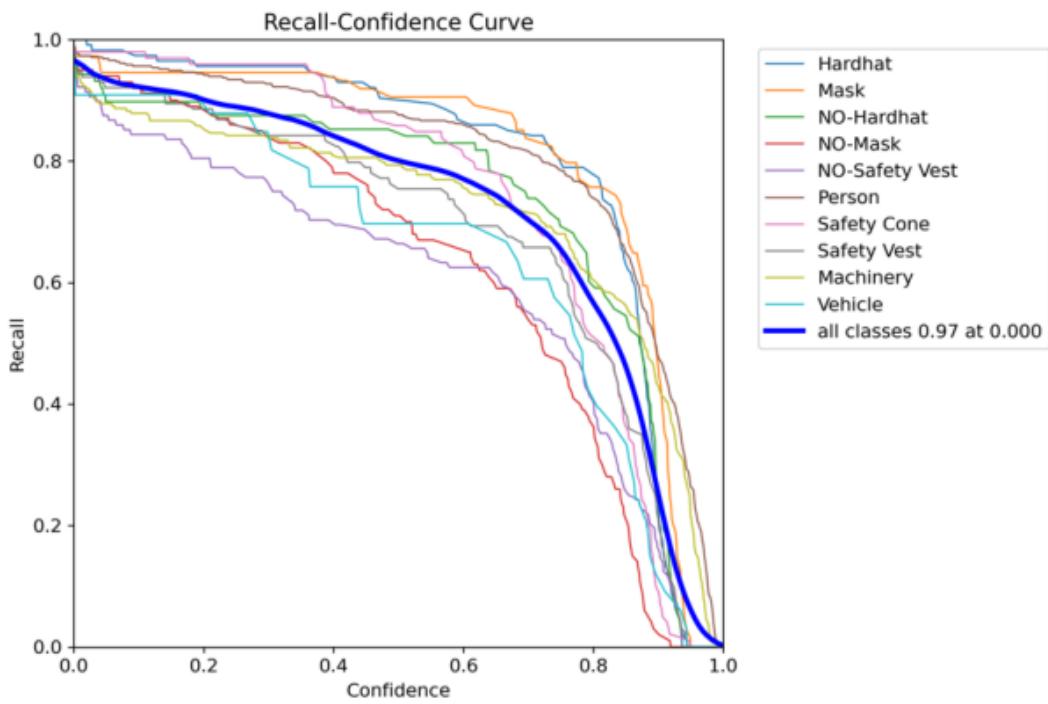
Visualization: P_curve.png



Visualization: results.png



Visualization: R_curve.png



2.0.18 Visualize prediction on test sets

```
[54]: # load the final trained model
model = YOLO("runs/detect/train4/weights/best.pt")

# Path to test images
test_images_dir = "datasets/images/test"

# get list of test images
test_images = [f for f in os.listdir(test_images_dir) if f.lower().endswith('.jpg', '.jpeg', '.png')]

# run inference and display predictions
# shuffle the images randomly
random.shuffle(test_images)

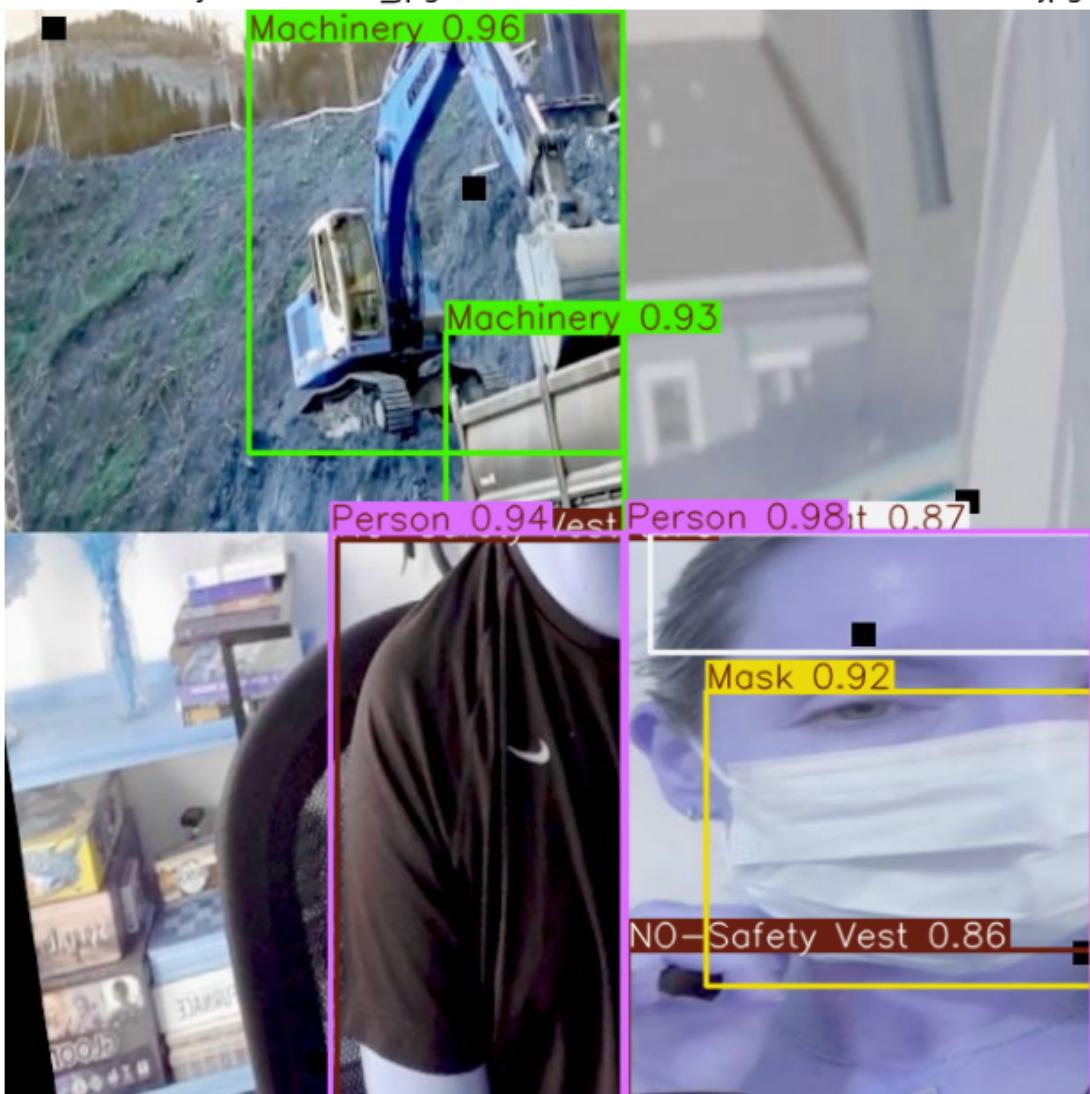
# limit to a subset
test_images_subset = test_images[:3]

# Run inference and display predictions for random images
for test_image in test_images_subset:
    image_path = os.path.join(test_images_dir, test_image)
    results = model.predict(source=image_path, save=False) # Run inference

    # Display the image with predictions
    img = results[0].plot() # YOLOv8's built-in method to plot predictions
    plt.figure(figsize=(10, 8))
    plt.imshow(img)
    plt.axis("off")
    plt.title(f"Predictions: {test_image}")
    plt.show()
```

```
image 1/1 C:\Users\DARiN\Documents\Python-JL\Github\aaI-501-final-project\notebooks\hybrid_model_train\datasets\images\test\youtube-497.jpg.rf.e0f0250ef7939eb4813e9b2f57b1a5ff.jpg: 640x640 1 Mask, 1 NO-Hardhat, 2 NO-Safety Vests, 2 Persons, 2 Machinerys, 267.9ms
Speed: 2.0ms preprocess, 267.9ms inference, 2.0ms postprocess per image at shape (1, 3, 640)
```

Predictions: youtube-497.jpg.rf.e0f0250ef7939eb4813e9b2f57b1a5ff.jpg



```
image 1/1 C:\Users\DAReN\Documents\Python-JL\Github\aaI-501-final-project\notebooks\hybrid_model_train\datasets\images\test\IMG_3103_mp4-15.jpg.rf.d7e8a84fc7e9d53fe348320ea3af0e94.jpg: 640x640 2 Hardhats, 1 Mask, 2 NO-Hardhats, 3 NO-Masks, 1 NO-Safety Vest, 6 Persons, 1 Safety Cone, 1 Machinery, 253.6ms
Speed: 2.0ms preprocess, 253.6ms inference, 1.0ms postprocess per image at shape (1, 3, 640, 640)
```

Predictions: IMG_3103_mp4-15.jpg.rf.d7e8a84fc7e9d53fe348320ea3af0e94.jpg

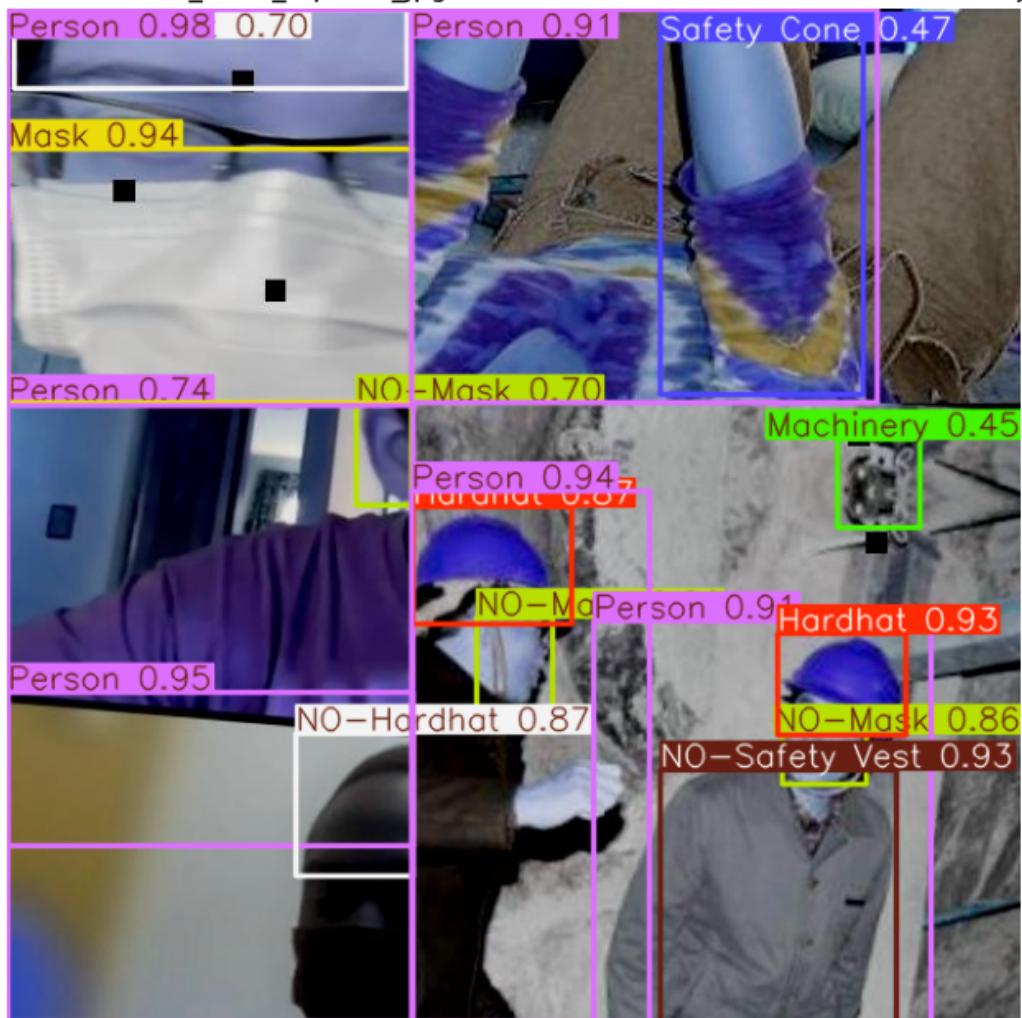
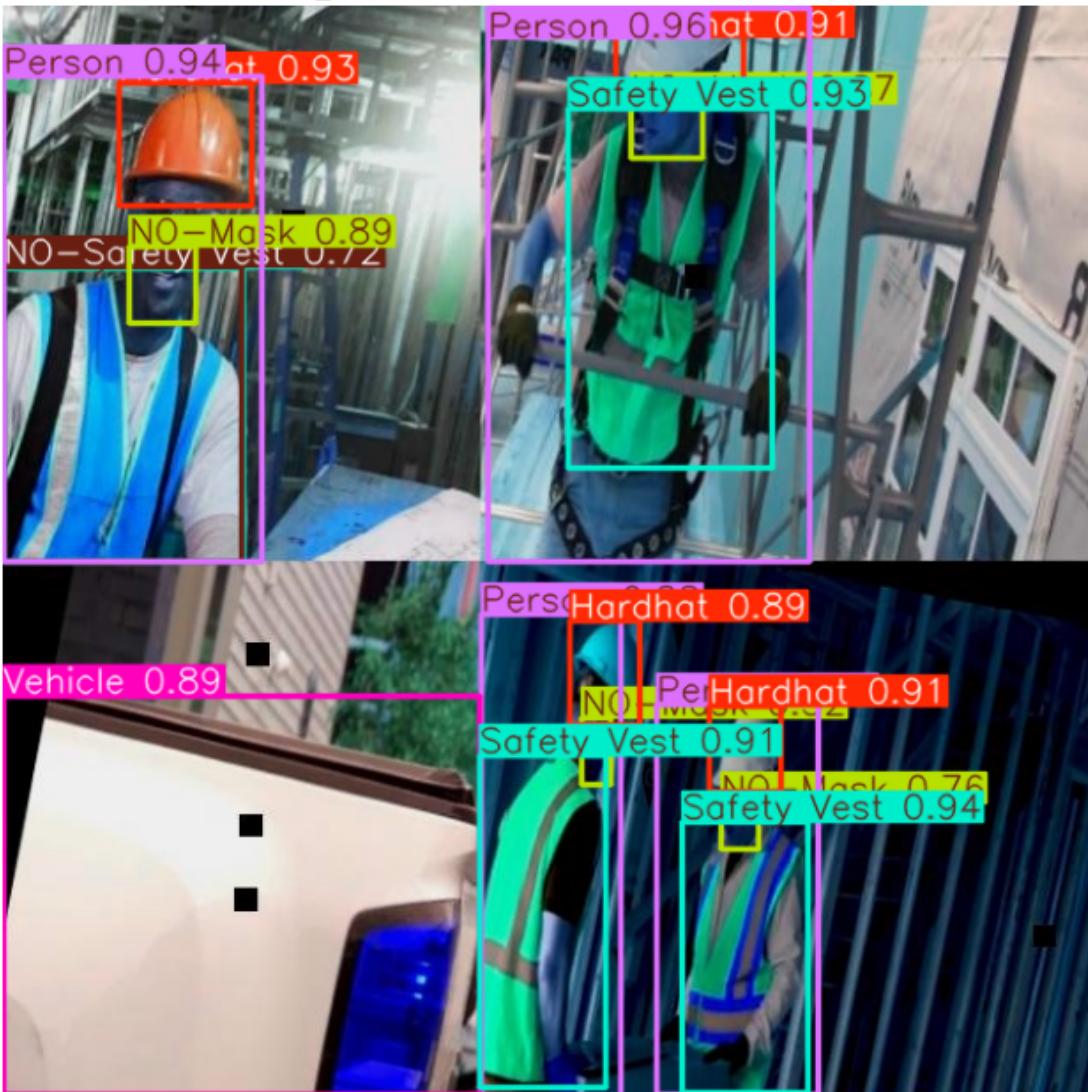


image 1/1 C:\Users\DARiN\Documents\Python-JL\Github\aaai-501-final-project\notebooks\hybrid_model_train\datasets\images\test\01261.jpg.rf.ee044d42291843f1358960c4d8c1d466.jpg: 640x640 4 Hardhats, 4 NO-Masks, 1 NO-Safety Vest, 4 Persons, 4 Safety Vests, 1 Vehicle, 235.2ms
Speed: 2.0ms preprocess, 235.2ms inference, 1.0ms postprocess per image at shape (1, 3, 640, 640)

Predictions: 01261.jpg.rf.ee044d42291843f1358960c4d8c1d466.jpg



2.0.19 Predict custom images

```
[11]: # Remove previous predictions
detect_path = "runs/detect/"
for folder in os.listdir(detect_path):
    folder_path = os.path.join(detect_path, folder)
    if os.path.isdir(folder_path) and folder.startswith("predict"):
        shutil.rmtree(folder_path, ignore_errors=True)

# Make predictions on "scene" images
model.predict()
```

```

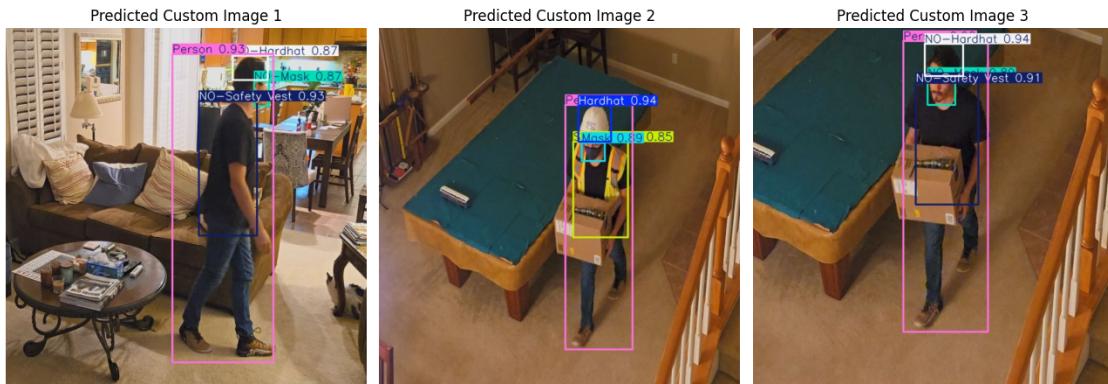
f"{test_images_dir}/scene*.jpg",
save=True,
conf=0.45,
agnostic_nms=True,
iou=0.8,
verbose=False
)

# Visualize the first three predicted "scene" images
scene_predicted_paths = glob.glob("runs/detect/predict/scene*.jpg")[:3]

if len(scene_predicted_paths) < 3:
    print("Not enough predictions found!")
else:
    fig, axes = plt.subplots(1, 3, figsize=(13, 6)) # 1 row, 3 columns
    for i, ax in enumerate(axes):
        image = mpimg.imread(scene_predicted_paths[i])
        ax.imshow(image)
        ax.axis('off')
        ax.set_title(f"Predicted Custom Image {i + 1}")
    plt.tight_layout()
    plt.show()

```

Results saved to runs\detect\predict



Model Prediction Comparison

December 5, 2024

1 Model Predictions Comparisons

AAI-501 Group 4: Fatimat Atanda, Victor Hugo Germano, Darin Verduzco

```
[1]: # import modules
import cv2 # opencv for images
import torch
from ultralytics import YOLO #for object detection
import matplotlib.pyplot as plt

# Support methods for groupd 4
import msaai
```

2 Load Model and Model Names

Using 4 main models:

- YOLO vanilla
- YOLO community pre trained on PPEs
- YOLO trained using manual annotations
- YOLO hybrid model

```
[3]: # load pre trained model
model_default = YOLO("../models/yolov8n.pt")
model_usd_g4 = YOLO("../models/best_custom.pt")
model_pre_trained = YOLO("../models/ppe.pt")
model_hybrid = YOLO("../models/hybrid_best.pt")

print(model_default.names)

print("===== USD Group 4 Class names =====")

print(model_usd_g4.names)

print(print("===== Community pre trained ====="))

print(model_pre_trained.names)
```

```
{0: 'person', 1: 'bicycle', 2: 'car', 3: 'motorcycle', 4: 'airplane', 5: 'bus',
6: 'train', 7: 'truck', 8: 'boat', 9: 'traffic light', 10: 'fire hydrant', 11:
```

```

'stop sign', 12: 'parking meter', 13: 'bench', 14: 'bird', 15: 'cat', 16: 'dog',
17: 'horse', 18: 'sheep', 19: 'cow', 20: 'elephant', 21: 'bear', 22: 'zebra',
23: 'giraffe', 24: 'backpack', 25: 'umbrella', 26: 'handbag', 27: 'tie', 28:
'suitcase', 29: 'frisbee', 30: 'skis', 31: 'snowboard', 32: 'sports ball', 33:
'kite', 34: 'baseball bat', 35: 'baseball glove', 36: 'skateboard', 37:
'surfboard', 38: 'tennis racket', 39: 'bottle', 40: 'wine glass', 41: 'cup', 42:
'fork', 43: 'knife', 44: 'spoon', 45: 'bowl', 46: 'banana', 47: 'apple', 48:
'sandwich', 49: 'orange', 50: 'broccoli', 51: 'carrot', 52: 'hot dog', 53:
'pizza', 54: 'donut', 55: 'cake', 56: 'chair', 57: 'couch', 58: 'potted plant',
59: 'bed', 60: 'dining table', 61: 'toilet', 62: 'tv', 63: 'laptop', 64:
'mouse', 65: 'remote', 66: 'keyboard', 67: 'cell phone', 68: 'microwave', 69:
'oven', 70: 'toaster', 71: 'sink', 72: 'refrigerator', 73: 'book', 74: 'clock',
75: 'vase', 76: 'scissors', 77: 'teddy bear', 78: 'hair drier', 79:
'toothbrush'}
===== USD Group 4 Class names =====
{0: 'Hardhat', 1: 'Mask', 2: 'NO-Hardhat', 3: 'NO-Mask', 4: 'NO-Safety Vest', 5:
'Person', 6: 'Safety Cone', 7: 'Safety Vest', 8: 'Machinery', 9: 'Vehicle'}
===== Community pre trained =====
None
{0: 'Hardhat', 1: 'Mask', 2: 'NO-Hardhat', 3: 'NO-Mask', 4: 'NO-Safety Vest', 5:
'Person', 6: 'Safety Cone', 7: 'Safety Vest', 8: 'machinery', 9: 'vehicle'}

```

```
[34]: # Predict the list of images and plot on a grid
def predict_ppes_grid(model, list_images, title=""):
    num_images = len(list_images)
    num_cols = min(num_images, 3) # minimum 2 columns
    num_rows = int(num_images / num_cols) + (1 if num_images % num_cols != 0
                                             else 0)

    # Create a grid of subplots.
    fig, axes = plt.subplots(num_rows, num_cols, figsize=(10, 9))
    list_axes = list(axes.flat)

    for i in range(num_images):
        file = list_images[i]
        # Open image file for reading and fix color grading
        img = cv2.imread(file)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        # perform inference using model and draw predictions
        results = model.predict(img, verbose=False)
        summary = msaaai.drawPredictions(results, img, model.names)
        list_axes[i].imshow(img)
        list_axes[i].axis('off')
        print(summary)

    plt.suptitle(title)
```

```
plt.tight_layout()  
plt.show()
```

3 Get predictions from models

Using the loaded models and multiple images as examples, understanding how the models performed

```
[35]: #Manually presenting annotations  
# train files  
base_dir="train_from_custom/datasets/raw_images/"  
  
image_files=[f"{base_dir}01437.jpg.rf.7839171efbe212f40b2c6ce1f3c8819b.jpg",  
            f"{base_dir}construction-619-.jpg.rf.  
            ↵8181765b21d618233610f2c91a6057a4.jpg",  
            f"{base_dir}construction-619-.jpg.rf.  
            ↵dc01af03cac6c0e921cf55d5d8e8465d.jpg",  
            f"{base_dir}005310.jpg.rf.83e6f0a246ede1e981765bf75cb7b067.jpg",  
            f"{base_dir}2008_008382.jpg.rf.88abdb29b7db6d6218399b4adebdbe66.  
            ↵jpg",  
            f"{base_dir}-2270-.png.jpg.rf.4b73cc556e91bc8c587c923e348717f9.jpg",  
            f"{base_dir}scene1_all_4.jpg",  
            f"{base_dir}scene2_all_3.jpg",  
            f"{base_dir}scene3_all_3.jpg",  
]  
  
predict_ppes_grid(model_default, image_files, title="YOLO vanilla model")  
  
1 person, 1 truck, 1 boat,  
6 persons,  
1 truck,  
2 persons, 1 bicycle,  
1 person, 1 train,  
4 persons, 1 bus,  
2 persons, 3 cups, 3 couchs, 1 dining table, 2 remotes, 1 book,  
1 person, 1 chair,  
2 persons, 1 bench, 2 cups, 1 chair, 1 dining table,
```

YOLO vanilla model



3.1 Pre-trained PEE Benchmark

```
[36]: predict_ppes_grid(model_pre_trained, image_files, title="YOLO Pre Trained Model")
```

3 Hardhats, 1 NO-Mask, 1 NO-Safety Vest, 1 Person, 2 machineries,
 2 Hardhats, 1 Mask, 1 NO-Hardhat, 1 NO-Mask, 3 NO-Safety Vests, 4 Persons, 4
 machineries,
 3 Persons, 5 machineries,
 2 Hardhats, 3 Persons, 2 Safety Vests, 7 machineries,
 1 Hardhat, 1 NO-Hardhat, 2 NO-Safety Vests, 2 Persons, 1 machinery,
 1 Hardhat, 1 Mask, 1 NO-Hardhat, 1 NO-Mask, 3 NO-Safety Vests, 3 Persons, 1
 Safety Cone, 1 machinery, 1 vehicle,
 1 Hardhat, 2 NO-Masks, 1 NO-Safety Vest, 1 Person,
 1 Hardhat, 1 Person, 1 Safety Vest,

1 Hardhat, 1 NO-Mask, 2 Persons, 2 Safety Vests,

YOLO Pre Trained model



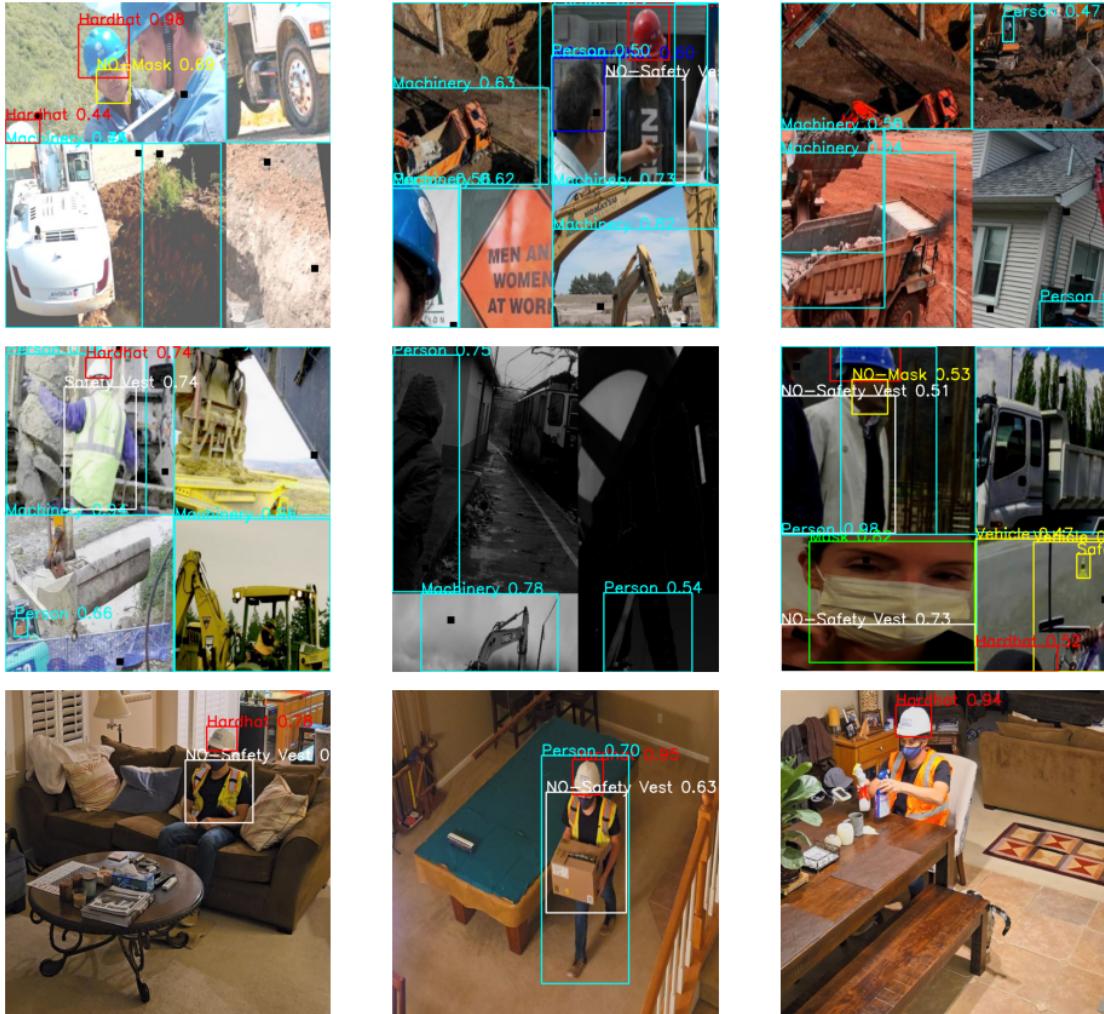
3.2 USD Group 4: Trained From Custom Dataset

```
[37]: predict_ppes_grid(model_usd_g4, image_files, title="YOLO USD Group 4: Trained  
From Custom Dataset")
```

2 Hardhats, 1 NO-Mask, 1 Person, 4 Machineries,
1 Hardhat, 1 NO-Hardhat, 1 NO-Mask, 1 NO-Safety Vest, 4 Persons, 6 Machineries,
1 Hardhat, 2 Persons, 6 Machineries,
1 Hardhat, 2 Persons, 1 Safety Vest, 4 Machineries,
1 NO-Safety Vest, 2 Persons, 1 Machinery,
2 Hardhats, 1 Mask, 3 NO-Masks, 2 NO-Safety Vests, 4 Persons, 1 Safety Cone, 1
Machinery, 2 Vehicles,
1 Hardhat, 1 NO-Safety Vest,

1 Hardhat, 1 NO-Safety Vest, 1 Person,
 1 Hardhat,

YOLO USD Group 4: Trained From Custom Dataset



3.3 USD Group 4: Hybrid Model

```
[38]: predict_ppes_grid(model_hybrid, image_files, title="YOLO USD Group 4: Hybrid Model")
```

2 Hardhats, 1 Person, 2 Machinerys,
 1 Hardhat, 2 NO-Hardhats, 1 NO-Mask, 1 NO-Safety Vest, 4 Persons, 3 Machinerys,
 1 Person, 6 Machinerys,
 1 Hardhat, 1 Person, 1 Safety Vest,
 1 NO-Hardhat, 2 Persons, 2 Machinerys,
 1 Hardhat, 1 Mask, 1 NO-Hardhat, 1 NO-Mask, 3 NO-Safety Vests, 3 Persons, 1
 Safety Cone, 1 Machinery, 1 Vehicle,

1 Hardhat, 1 Person,
1 Hardhat, 1 Mask, 1 Person, 1 Safety Vest,
1 Hardhat, 1 NO-Mask, 1 Person,

YOLO USD Group 4: Hybrid Model



Video Stream Prediction

December 5, 2024

1 Video Recognition exercise

AAI-501 Group 4: Fatimat Atanda, Victor Hugo Germano, Darin Verduzco

1.1 Configuring the environment

```
[1]: import cv2
from ultralytics import YOLO
import math
import os # file

import msaaai

# MODEL loading and config
# model
model = YOLO("../models/best_custom.pt")

class_names = ['Hardhat', 'Mask', 'NO-Hardhat',
               'NO-Mask', 'NO-Safety Vest', 'Person',
               'Safety Cone', 'Safety Vest', 'Machinery', 'Vehicle']

# Create output folder
os.makedirs("../datasets/stream_result", exist_ok=True)
```

```
[2]: # DRAWING METHODS

def predict_video_stream(camera, output, verbose=False, showWindow=True, flipImage=False):
    while True:
        ret, frame = camera.read()
        if not ret:
            break

        if flipImage:
            frame = cv2.flip(frame, 1)

    results = model(frame, stream=True, verbose=False)
    msaaai.drawPredictions(results, frame, class_names)
```

```

# Write the frame to the output file
output.write(frame)
# Display the captured frame - Performance changes when not showing
if showWindow:
    cv2.imshow('Camera', frame)

# Press 'q' to exit the loop
if cv2.waitKey(1) == ord('q'):
    break

```

1.2 Capture Livestream Video from webcam

1.2.1 Use 'q' to close the window

[3]:

```

# Open the default camera
cam = cv2.VideoCapture(0)

# Get the default frame width and height
frame_width = int(cam.get(cv2.CAP_PROP_FRAME_WIDTH))
frame_height = int(cam.get(cv2.CAP_PROP_FRAME_HEIGHT))

# Define the codec and create VideoWriter object
fourcc = cv2.VideoWriter_fourcc(*'mp4v')
out = cv2.VideoWriter('../dataset/stream_result/output_webcam.mp4', fourcc, 20.
                     ↴0, (frame_width, frame_height))

# START PREDICTION AND PRESENTATION
predict_video_stream(cam, out, verbose=False, showWindow=True, flipImage=True)

# Release the capture and writer objects
cam.release()
out.release()
cv2.destroyAllWindows()
cv2.waitKey(1)

```

2024-11-28 18:35:26.848 python3[53673:8311993] WARNING:
AVCaptureDeviceTypeExternal is deprecated for Continuity Cameras. Please use
AVCaptureDeviceTypeContinuityCamera and add
NSCameraUseContinuityCameraDeviceType to your Info.plist.

[3]: -1

2 Analyzing Video File

Video download link:

https://uofsandiego-my.sharepoint.com/:v/g/personal/dverduzco_sandiego_edu/Ebp8SRL0ZI9F

```
[3]: # Open the default camera
cam = cv2.VideoCapture("../datasets/custom_ppe_videos/scene1.mp4")

# Get the default frame width and height
frame_width = int(cam.get(cv2.CAP_PROP_FRAME_WIDTH))
frame_height = int(cam.get(cv2.CAP_PROP_FRAME_HEIGHT))

# Define the codec and create VideoWriter object
fourcc = cv2.VideoWriter_fourcc(*'mp4v')
out = cv2.VideoWriter('../datasets/stream_result/output_custom_video.mp4', fourcc, 20.0, (frame_width, frame_height))

# START PREDICTION AND PRESENTATION
predict_video_stream(cam, out, verbose=False, showWindow=True)

# Release the capture and writer objects
cam.release()
out.release()
cv2.destroyAllWindows()
cv2.waitKey(1)
```

```
2024-11-28 18:51:47.362 python3[54299:8328625] +[IMKClient subclass]: chose IMKClient_Modern
2024-11-28 18:51:47.362 python3[54299:8328625] +[IMKInputSession subclass]: chose IMKInputSession_Modern
```

```
[3]: -1
```